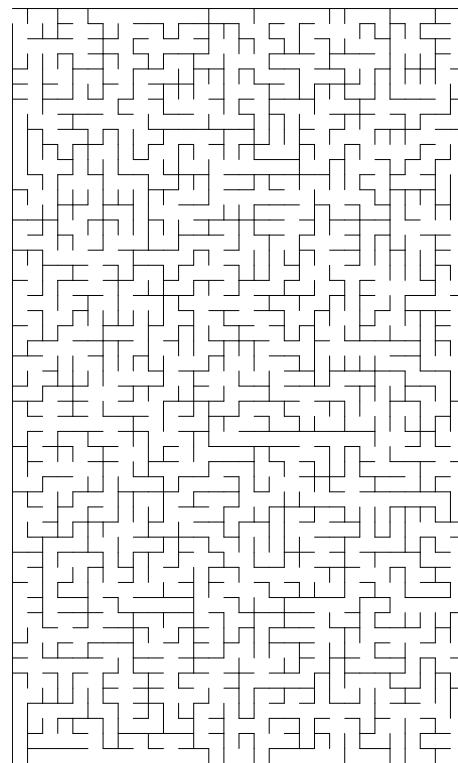


Algorithms and Analysis

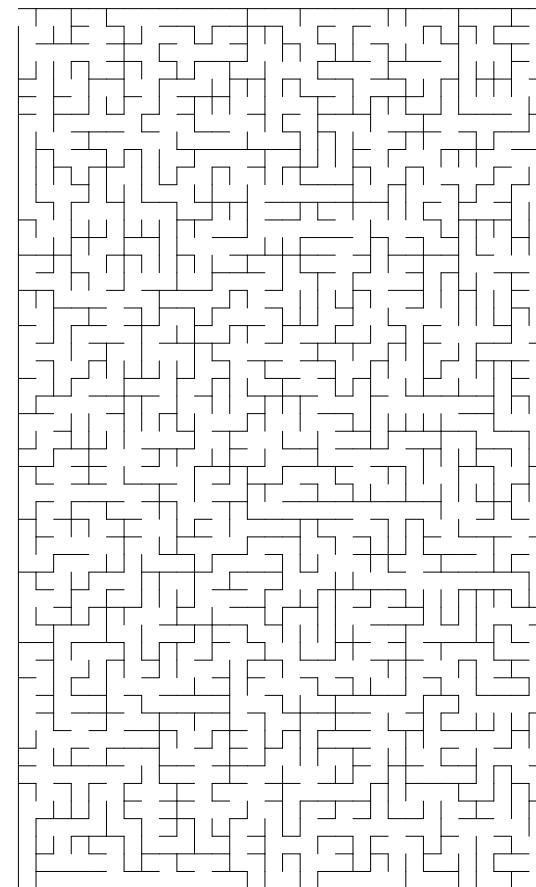
Lesson 14: *Use Arrays for Fast Set Algorithms*



Equivalent classes, Disjoint Set, Fast Sets

Outline

1. Equivalent Classes
2. Disjoint Sets
3. Fast Sets



Equivalence Relations

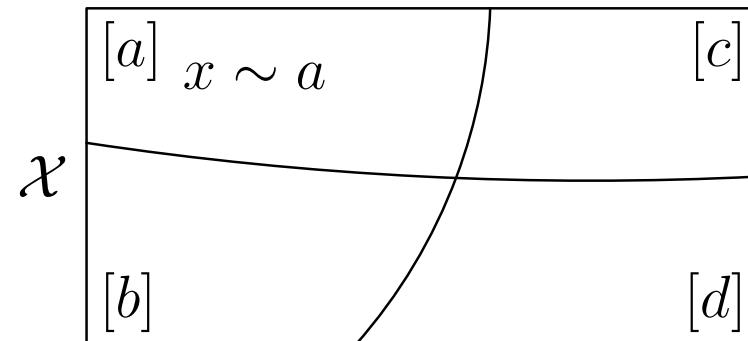
- Given a set of elements $\mathcal{X} = \{x_1, x_2, \dots\}$ and a binary relationship \sim with the following properties

(Reflexivity) For every element $x \in \mathcal{X}$, $x \sim x$

(Symmetry) For every two elements $x, y \in \mathcal{X}$ if $x \sim y$ then $y \sim x$

(Transitivity) For every three elements $x, y, z \in \mathcal{X}$ if $x \sim y$ and $y \sim z$ then $x \sim z$

- Then \sim defines a partitioning of the set into **equivalence classes**



Equivalence Relations

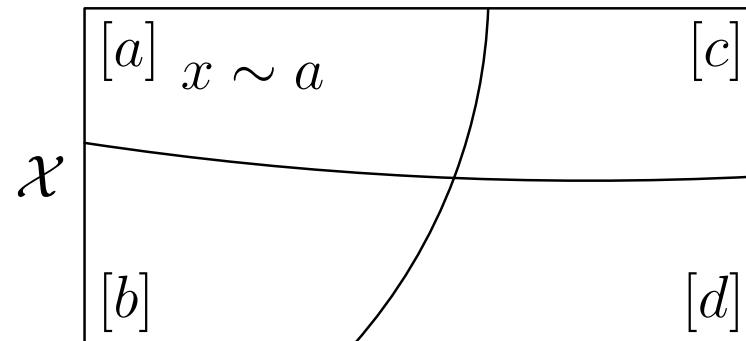
- Given a set of elements $\mathcal{X} = \{x_1, x_2, \dots\}$ and a binary relationship \sim with the following properties

(Reflexivity) For every element $x \in \mathcal{X}$, $x \sim x$

(Symmetry) For every two elements $x, y \in \mathcal{X}$ if $x \sim y$ then $y \sim x$

(Transitivity) For every three elements $x, y, z \in \mathcal{X}$ if $x \sim y$ and $y \sim z$ then $x \sim z$

- Then \sim defines a partitioning of the set into **equivalence classes**



Equivalence Relations

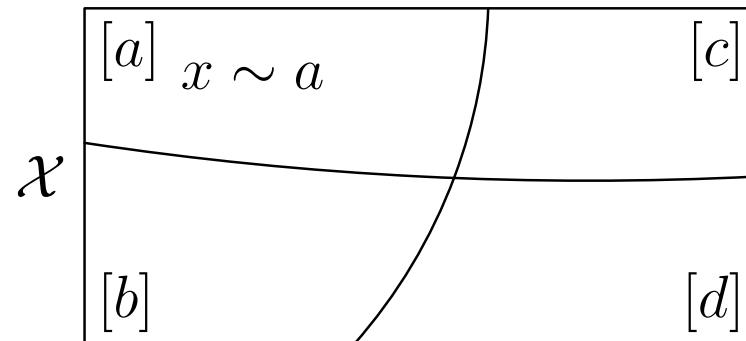
- Given a set of elements $\mathcal{X} = \{x_1, x_2, \dots\}$ and a binary relationship \sim with the following properties

(Reflexivity) For every element $x \in \mathcal{X}$, $x \sim x$

(Symmetry) For every two elements $x, y \in \mathcal{X}$ if $x \sim y$ then
 $y \sim x$

(Transitivity) For every three elements $x, y, z \in \mathcal{X}$ if $x \sim y$ and $y \sim z$ then $x \sim z$

- Then \sim defines a partitioning of the set into **equivalence classes**



Equivalence Relations

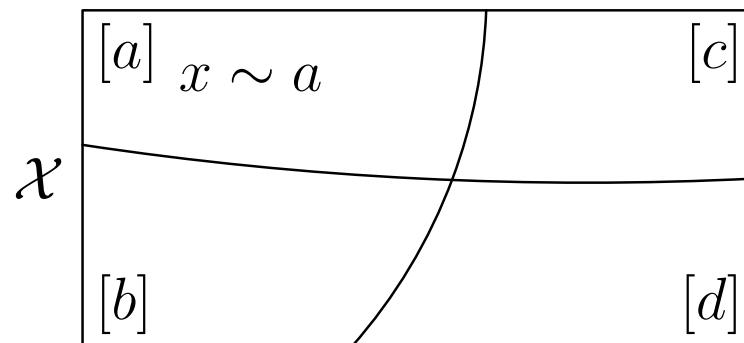
- Given a set of elements $\mathcal{X} = \{x_1, x_2, \dots\}$ and a binary relationship \sim with the following properties

(Reflexivity) For every element $x \in \mathcal{X}$, $x \sim x$

(Symmetry) For every two elements $x, y \in \mathcal{X}$ if $x \sim y$ then $y \sim x$

(Transitivity) For every three elements $x, y, z \in \mathcal{X}$ if $x \sim y$ and $y \sim z$ then $x \sim z$

- Then \sim defines a partitioning of the set into **equivalence classes**



Equivalence Relations

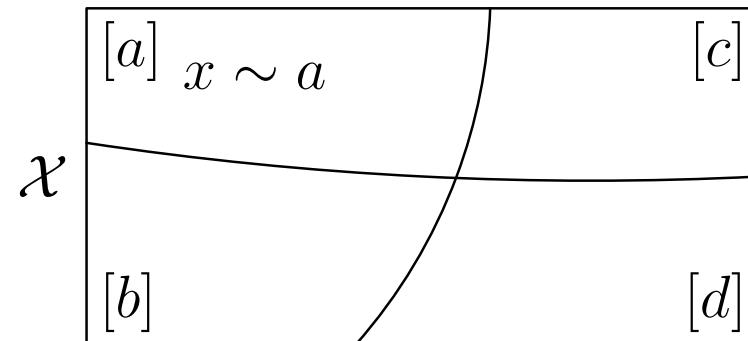
- Given a set of elements $\mathcal{X} = \{x_1, x_2, \dots\}$ and a binary relationship \sim with the following properties

(Reflexivity) For every element $x \in \mathcal{X}$, $x \sim x$

(Symmetry) For every two elements $x, y \in \mathcal{X}$ if $x \sim y$ then $y \sim x$

(Transitivity) For every three elements $x, y, z \in \mathcal{X}$ if $x \sim y$ and $y \sim z$ then $x \sim z$

- Then \sim defines a partitioning of the set into **equivalence classes**



Example of Equivalence Classes

- Although, equivalent classes sound very mathematical they often provide a useful formalisation of the real world
- E.g. Pairs of web pages with a link in each direction between them
- Consider web pages in the same equivalence class if you can get from one to the other by clicking links
- Partitions the web into linked domains
- Friendship relations in social media

Example of Equivalence Classes

- Although, equivalent classes sound very mathematical they often provide a useful formalisation of the real world
- E.g. Pairs of web pages with a link in each direction between them
- Consider web pages in the same equivalence class if you can get from one to the other by clicking links
- Partitions the web into linked domains
- Friendship relations in social media

Example of Equivalence Classes

- Although, equivalent classes sound very mathematical they often provide a useful formalisation of the real world
- E.g. Pairs of web pages with a link in each direction between them
- Consider web pages in the same equivalence class if you can get from one to the other by clicking links
- Partitions the web into linked domains
- Friendship relations in social media

Example of Equivalence Classes

- Although, equivalent classes sound very mathematical they often provide a useful formalisation of the real world
- E.g. Pairs of web pages with a link in each direction between them
- Consider web pages in the same equivalence class if you can get from one to the other by clicking links
- Partitions the web into linked domains
- Friendship relations in social media

Example of Equivalence Classes

- Although, equivalent classes sound very mathematical they often provide a useful formalisation of the real world
- E.g. Pairs of web pages with a link in each direction between them
- Consider web pages in the same equivalence class if you can get from one to the other by clicking links
- Partitions the web into linked domains
- Friendship relations in social media

Dynamic Equivalence Classes

- Finding equivalence classes is rather easy using graph traversal algorithms
- However, as our web example suggests, there are applications where equivalence classes change over time
- Adding a link could join two domains which were separate
- We will see this is a useful idea both for building mazes and (in a later lecture) for finding minimum spanning trees
- Building a data structure which finds equivalence classes where the equivalence relation changes over time is challenging

Dynamic Equivalence Classes

- Finding equivalence classes is rather easy using graph traversal algorithms
- However, as our web example suggests, there are applications where equivalence classes change over time
- Adding a link could join two domains which were separate
- We will see this is a useful idea both for building mazes and (in a later lecture) for finding minimum spanning trees
- Building a data structure which finds equivalence classes where the equivalence relation changes over time is challenging

Dynamic Equivalence Classes

- Finding equivalence classes is rather easy using graph traversal algorithms
- However, as our web example suggests, there are applications where equivalence classes change over time
- Adding a link could join two domains which were separate
- We will see this is a useful idea both for building mazes and (in a later lecture) for finding minimum spanning trees
- Building a data structure which finds equivalence classes where the equivalence relation changes over time is challenging

Dynamic Equivalence Classes

- Finding equivalence classes is rather easy using graph traversal algorithms
- However, as our web example suggests, there are applications where equivalence classes change over time
- Adding a link could join two domains which were separate
- We will see this is a useful idea both for building mazes and (in a later lecture) for finding minimum spanning trees
- Building a data structure which finds equivalence classes where the equivalence relation changes over time is challenging

Dynamic Equivalence Classes

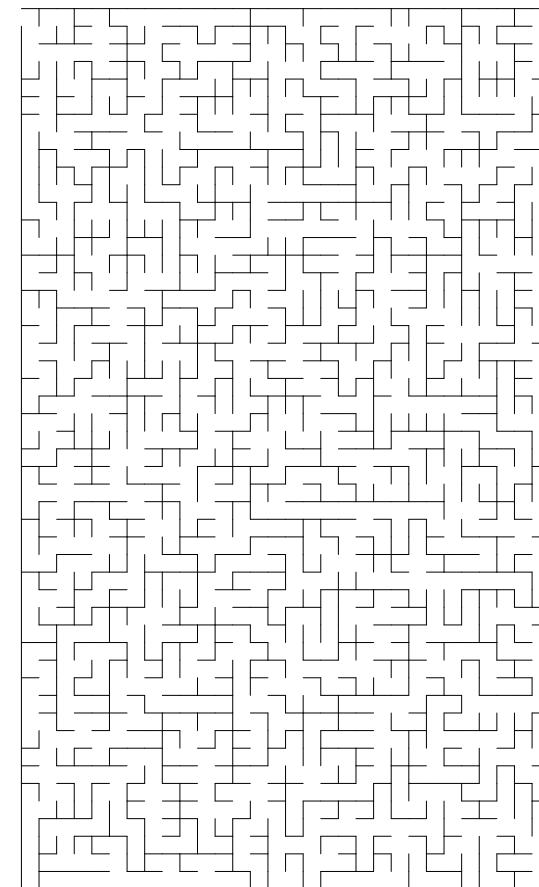
- Finding equivalence classes is rather easy using graph traversal algorithms
- However, as our web example suggests, there are applications where equivalence classes change over time
- Adding a link could join two domains which were separate
- We will see this is a useful idea both for building mazes and (in a later lecture) for finding minimum spanning trees
- Building a data structure which finds equivalence classes where the equivalence relation changes over time is challenging

Dynamic Equivalence Classes

- Finding equivalence classes is rather easy using graph traversal algorithms
- However, as our web example suggests, there are applications where equivalence classes change over time
- Adding a link could join two domains which were separate
- We will see this is a useful idea both for building mazes and (in a later lecture) for finding minimum spanning trees
- Building a data structure which finds equivalence classes where the equivalence relation changes over time is challenging, **but fortunately there is an elegant solution to this**

Outline

1. Equivalent Classes
2. Disjoint Sets
3. Fast Sets



Union-Find

- In the union-find algorithm we have a set of objects $x \in \mathcal{S}$ which are to be grouped into equivalence classes $\mathcal{S}_1, \mathcal{S}_2, \dots$
- Initially each object is in its own equivalence class
- **union** combines two equivalence classes
- **find** returns a unique member of the equivalence class (e.g $[a]$)
- This is a common problem for which we will write a class `DisjointSets` to perform fast unions and finds

Union-Find

- In the union-find algorithm we have a set of objects $x \in \mathcal{S}$ which are to be grouped into equivalence classes $\mathcal{S}_1, \mathcal{S}_2, \dots$
- Initially each object is in its own equivalence class
- **union** combines two equivalence classes
- **find** returns a unique member of the equivalence class (e.g $[a]$)
- This is a common problem for which we will write a class `DisjointSets` to perform fast unions and finds

Union-Find

- In the union-find algorithm we have a set of objects $x \in \mathcal{S}$ which are to be grouped into equivalence classes $\mathcal{S}_1, \mathcal{S}_2, \dots$
- Initially each object is in its own equivalence class
- **union** combines two equivalence classes
- **find** returns a unique member of the equivalence class (e.g $[a]$)
- This is a common problem for which we will write a class `DisjointSets` to perform fast unions and finds

Union-Find

- In the union-find algorithm we have a set of objects $x \in \mathcal{S}$ which are to be grouped into equivalence classes $\mathcal{S}_1, \mathcal{S}_2, \dots$
- Initially each object is in its own equivalence class
- **union** combines two equivalence classes
- **find** returns a unique member of the equivalence class (e.g $[a]$)
- This is a common problem for which we will write a class `DisjointSets` to perform fast unions and finds

Union-Find

- In the union-find algorithm we have a set of objects $x \in \mathcal{S}$ which are to be grouped into equivalence classes $\mathcal{S}_1, \mathcal{S}_2, \dots$
- Initially each object is in its own equivalence class
- **union** combines two equivalence classes
- **find** returns a unique member of the equivalence class (e.g $[a]$)
- This is a common problem for which we will write a class `DisjointSets` to perform fast unions and finds

DisjointSets

- We want to create a class

```
class DisjointSets
{
    DisjointSets(int numElements) /* Constructor */
    int find(int x) /* Find root */
    void union_(int root1, int root2) /* Union */

private:
    int* s;
}
```

- Where `find(x)` returns a unique identifier for the subset which element `x` belongs to
- The array `s` contains labelling information to implement `find(x)`

DisjointSets

- We want to create a class

```
class DisjointSets
{
    DisjointSets(int numElements) /* Constructor */
    int find(int x) /* Find root */
    void union_(int root1, int root2) /* Union */

    private:
        int* s;
}
```

- Where `find(x)` returns a unique identifier for the subset which element `x` belongs to
- The array `s` contains labelling information to implement `find(x)`

DisjointSets

- We want to create a class

```
class DisjointSets
{
    DisjointSets(int numElements) /* Constructor */
    int find(int x) /* Find root */
    void union_(int root1, int root2) /* Union */

    private:
        int* s;
}
```

- Where `find(x)` returns a unique identifier for the subset which element `x` belongs to
- The array `s` contains labelling information to implement `find(x)`

The Union-Find Dilemma

- A natural algorithm to perform finds is to maintain an array returning a subset label for each element—this makes find fast
- However, every time we combine two subset we have to change all the labels in this array (taking $O(n)$ operations)
- If we are unlucky the cost of performing n unions is $\Theta(n^2)$
- If we ensure that we relabel the smaller subset then the time complexity is $\Theta(n \log(n))$
- Fast *finds* seems to give slow(ish) *unions*
- What about the other way around?

The Union-Find Dilemma

- A natural algorithm to perform finds is to maintain an array returning a subset label for each element—**this makes find fast**
- However, every time we combine two subset we have to change all the labels in this array (taking $O(n)$ operations)
- If we are unlucky the cost of performing n unions is $\Theta(n^2)$
- If we ensure that we relabel the smaller subset then the time complexity is $\Theta(n \log(n))$
- Fast *finds* seems to give slow(ish) *unions*
- What about the other way around?

The Union-Find Dilemma

- A natural algorithm to perform finds is to maintain an array returning a subset label for each element—this makes `find` fast
- However, every time we combine two subset we have to change all the labels in this array (taking $O(n)$ operations)
- If we are unlucky the cost of performing n unions is $\Theta(n^2)$
- If we ensure that we relabel the smaller subset then the time complexity is $\Theta(n \log(n))$
- Fast *finds* seems to give slow(ish) *unions*
- What about the other way around?

The Union-Find Dilemma

- A natural algorithm to perform finds is to maintain an array returning a subset label for each element—this makes `find` fast
- However, every time we combine two subset we have to change all the labels in this array (taking $O(n)$ operations)
- If we are unlucky the cost of performing n unions is $\Theta(n^2)$
- If we ensure that we relabel the smaller subset then the time complexity is $\Theta(n \log(n))$
- Fast *finds* seems to give slow(ish) *unions*
- What about the other way around?

The Union-Find Dilemma

- A natural algorithm to perform finds is to maintain an array returning a subset label for each element—this makes `find` fast
- However, every time we combine two subset we have to change all the labels in this array (taking $O(n)$ operations)
- If we are unlucky the cost of performing n unions is $\Theta(n^2)$
- If we ensure that we relabel the smaller subset then the time complexity is $\Theta(n \log(n))$
- Fast *finds* seems to give slow(ish) *unions*
- What about the other way around?

The Union-Find Dilemma

- A natural algorithm to perform finds is to maintain an array returning a subset label for each element—this makes `find` fast
- However, every time we combine two subset we have to change all the labels in this array (taking $O(n)$ operations)
- If we are unlucky the cost of performing n unions is $\Theta(n^2)$
- If we ensure that we relabel the smaller subset then the time complexity is $\Theta(n \log(n))$
- *Fast `finds` seems to give slow(ish) `unions`*
- What about the other way around?

The Union-Find Dilemma

- A natural algorithm to perform finds is to maintain an array returning a subset label for each element—this makes `find` fast
- However, every time we combine two subset we have to change all the labels in this array (taking $O(n)$ operations)
- If we are unlucky the cost of performing n unions is $\Theta(n^2)$
- If we ensure that we relabel the smaller subset then the time complexity is $\Theta(n \log(n))$
- Fast *finds* seems to give slow(ish) *unions*
- What about the other way around?

Fast Union

- To achieve fast unions we can represent our disjoint sets as a forest (many disjoint trees)
- Every time we perform a union we make one of the trees point to the head of the other tree
- The cost of `find` depends on the depth of the tree
- To make unions efficient we make the shallow tree a subtree of the deeper tree

Fast Union

- To achieve fast unions we can represent our disjoint sets as a forest (many disjoint trees)
- Every time we perform a union we make one of the trees point to the head of the other tree
- The cost of find depends on the depth of the tree
- To make unions efficient we make the shallow tree a subtree of the deeper tree

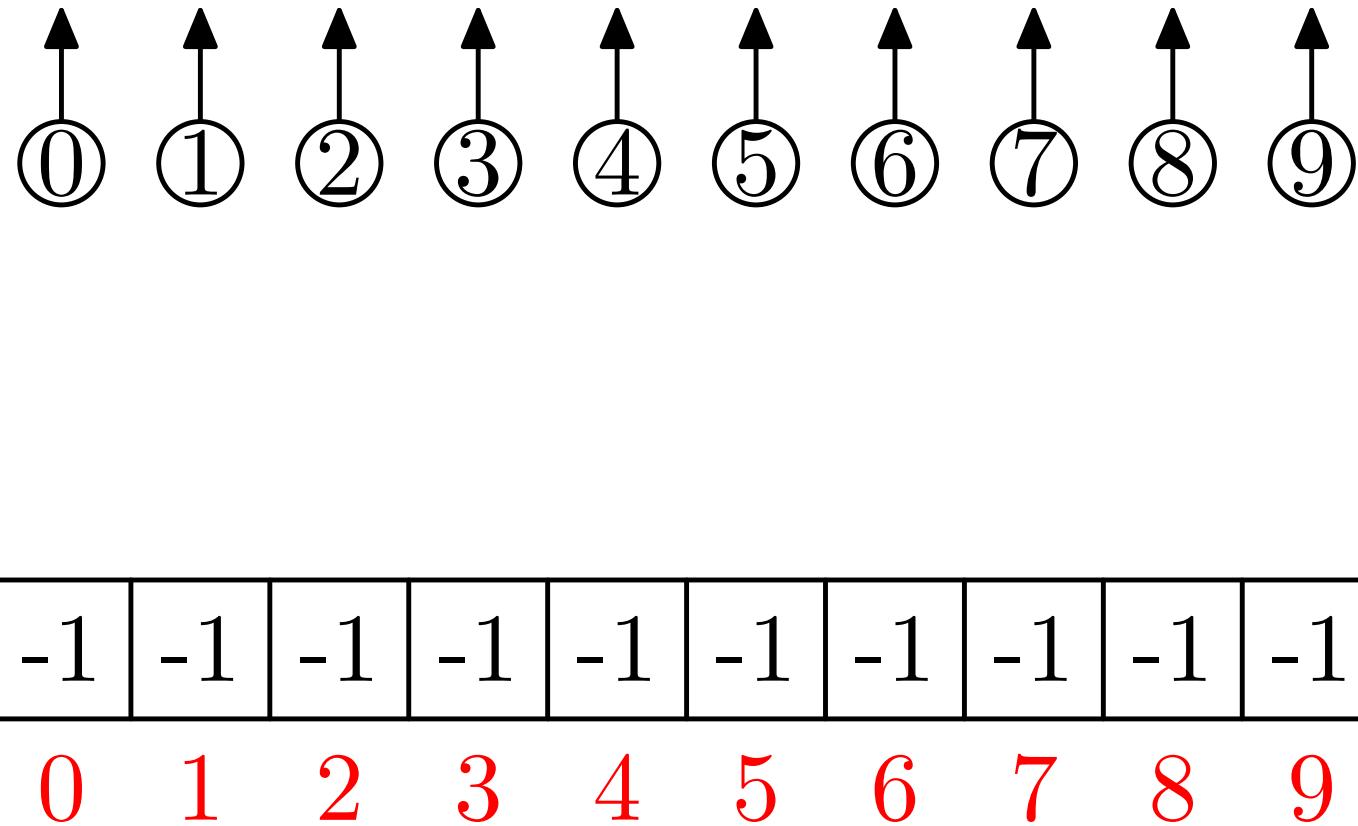
Fast Union

- To achieve fast unions we can represent our disjoint sets as a forest (many disjoint trees)
- Every time we perform a union we make one of the trees point to the head of the other tree
- The cost of `find` depends on the depth of the tree
- To make unions efficient we make the shallow tree a subtree of the deeper tree

Fast Union

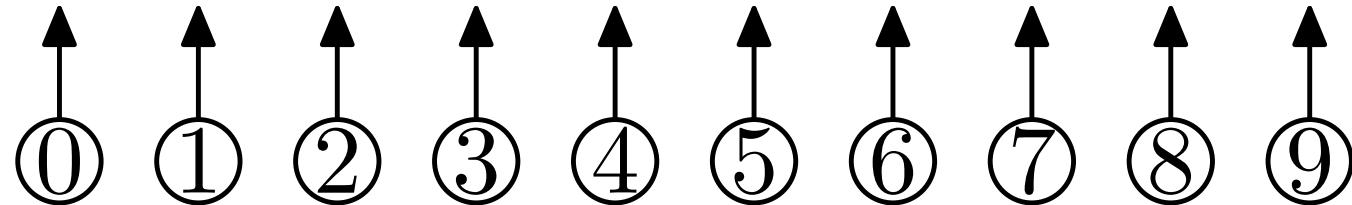
- To achieve fast unions we can represent our disjoint sets as a forest (many disjoint trees)
- Every time we perform a union we make one of the trees point to the head of the other tree
- The cost of `find` depends on the depth of the tree
- To make unions efficient we make the shallow tree a subtree of the deeper tree

Putting it Together



Putting it Together

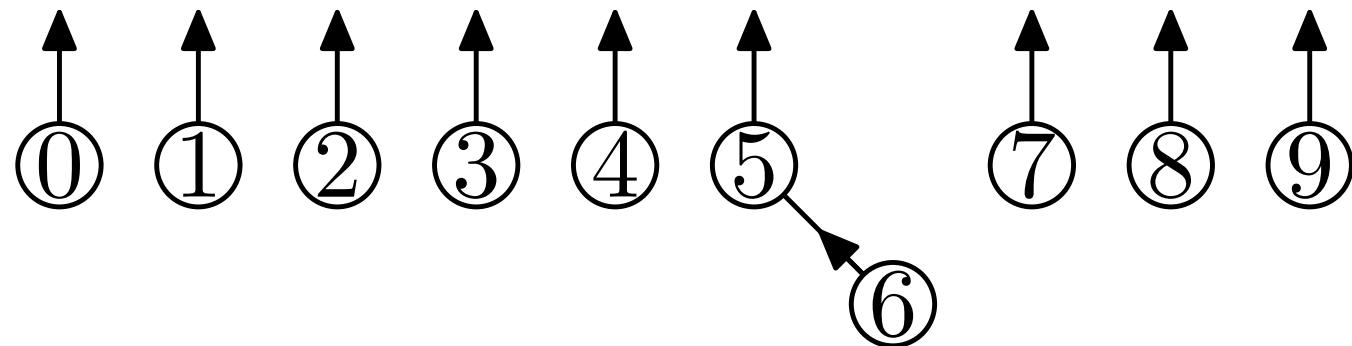
`union(find(5),find(6))`



-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

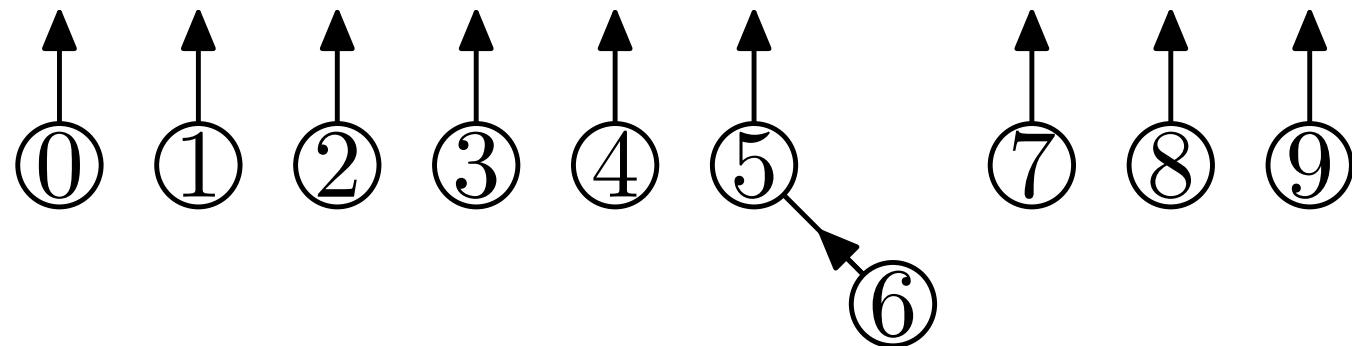
`union(find(5),find(6))`



-1	-1	-1	-1	-1	-2	5	-1	-1	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

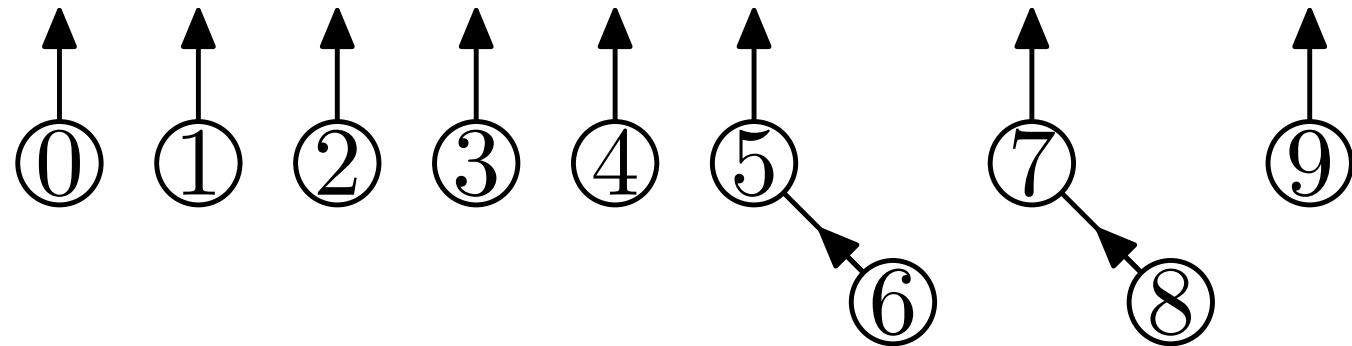
`union(find(7),find(8))`



-1	-1	-1	-1	-1	-2	5	-1	-1	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

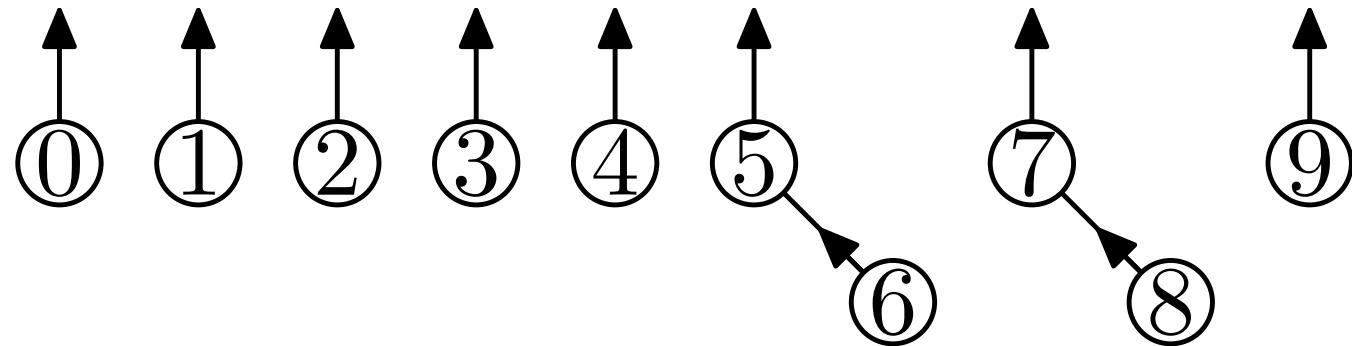
`union(find(7),find(8))`



-1	-1	-1	-1	-1	-2	5	-2	7	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

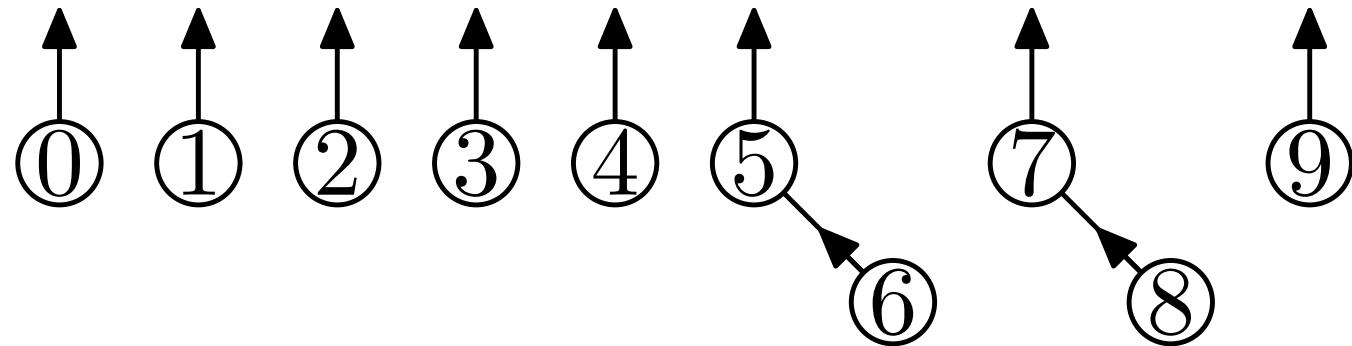
`find(6)`



-1	-1	-1	-1	-1	-2	5	-2	7	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

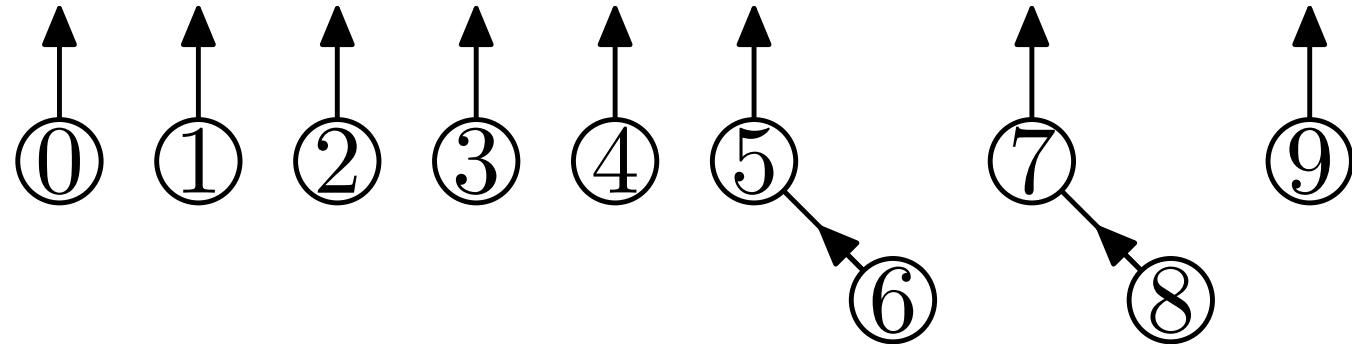
$\text{find}(6)=5$



-1	-1	-1	-1	-1	-2	5	-2	7	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

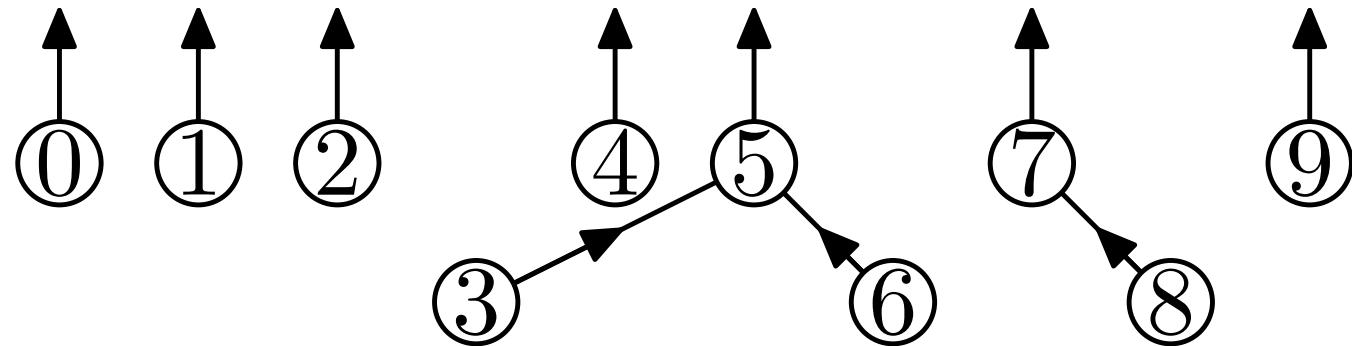
`union(find(3),find(6))`



-1	-1	-1	-1	-1	-2	5	-2	7	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

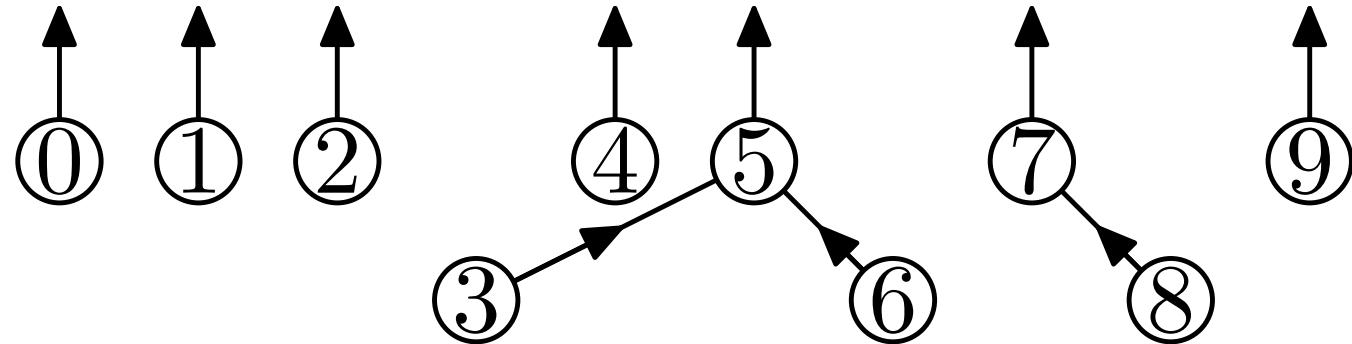
`union(find(3),find(6))`



-1	-1	-1	5	-1	-2	5	-2	7	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

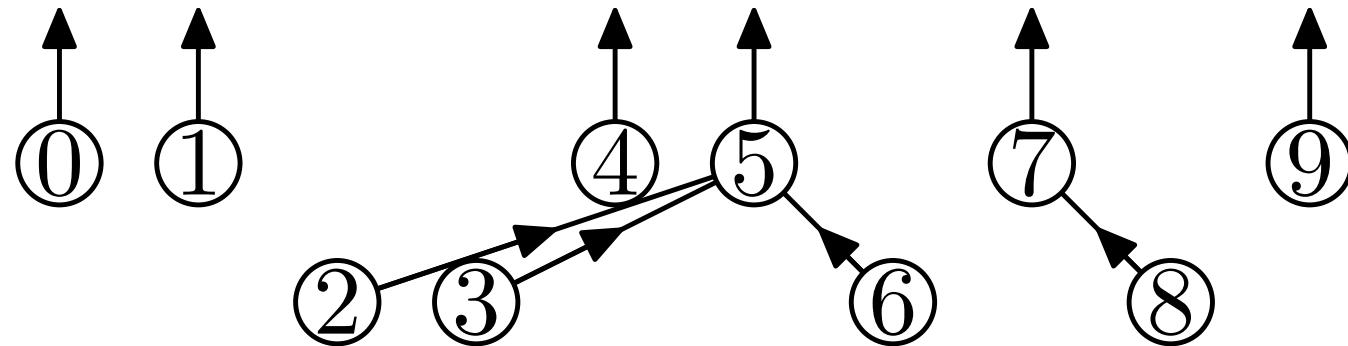
`union(find(2),find(6))`



-1	-1	-1	5	-1	-2	5	-2	7	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

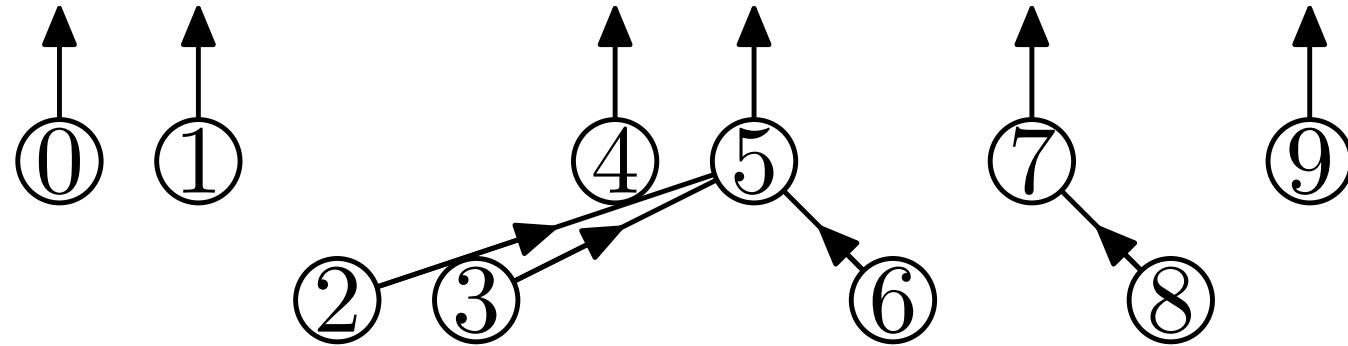
`union(find(2),find(6))`



-1	-1	5	5	-1	-2	5	-2	7	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

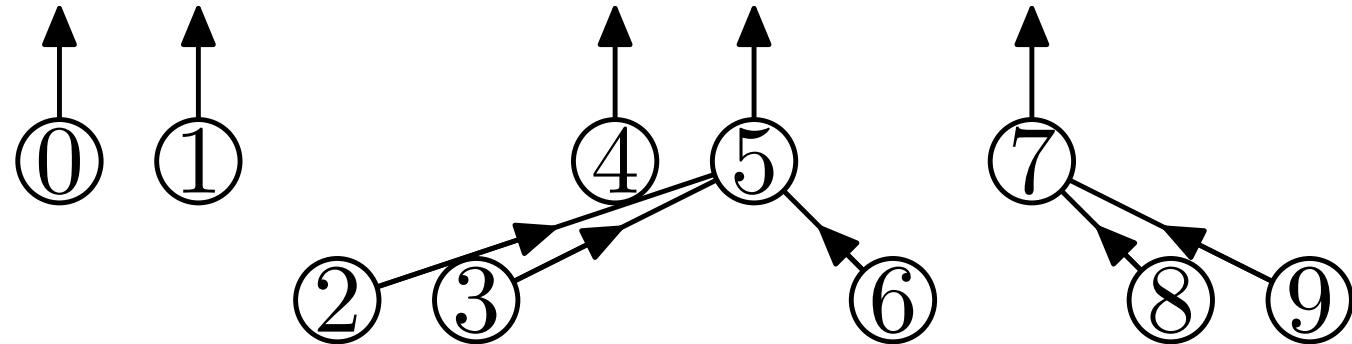
`union(find(9),find(8))`



-1	-1	5	5	-1	-2	5	-2	7	-1
0	1	2	3	4	5	6	7	8	9

Putting it Together

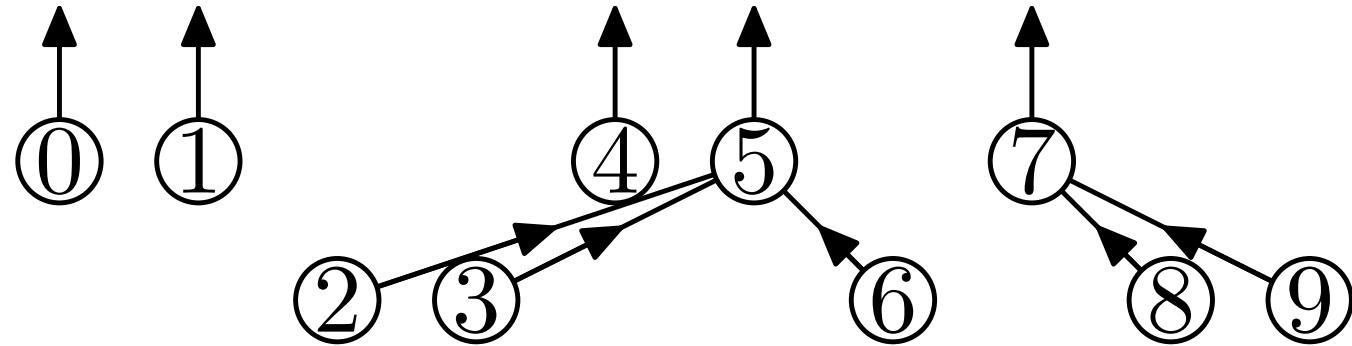
`union(find(9),find(8))`



-1	-1	5	5	-1	-2	5	-2	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

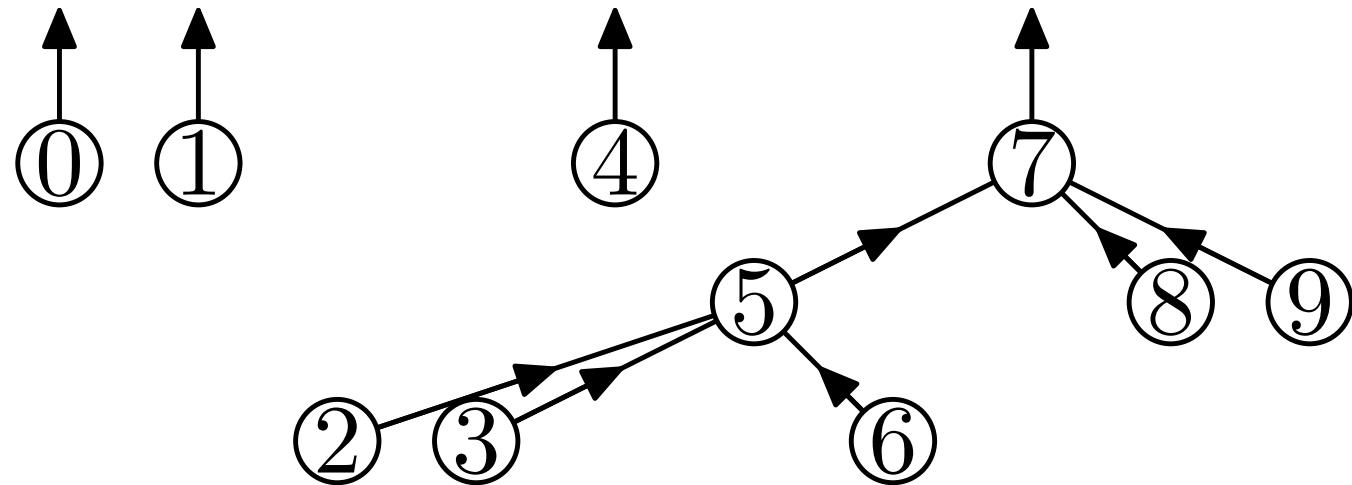
`union(find(9),find(3))`



-1	-1	5	5	-1	-2	5	-2	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

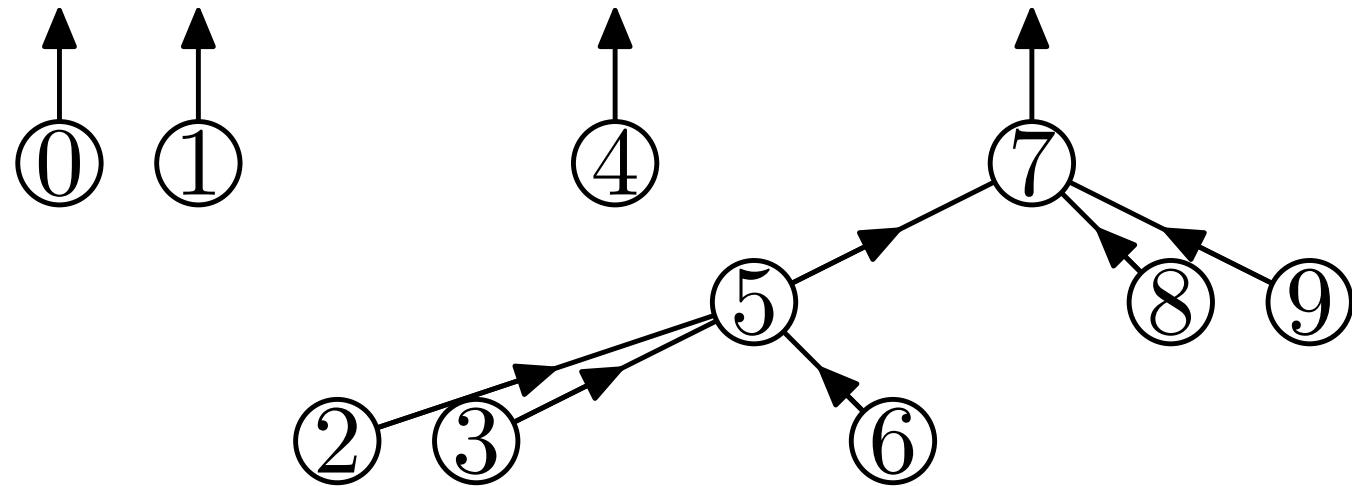
`union(find(9),find(3))`



-1	-1	5	5	-1	7	5	-3	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

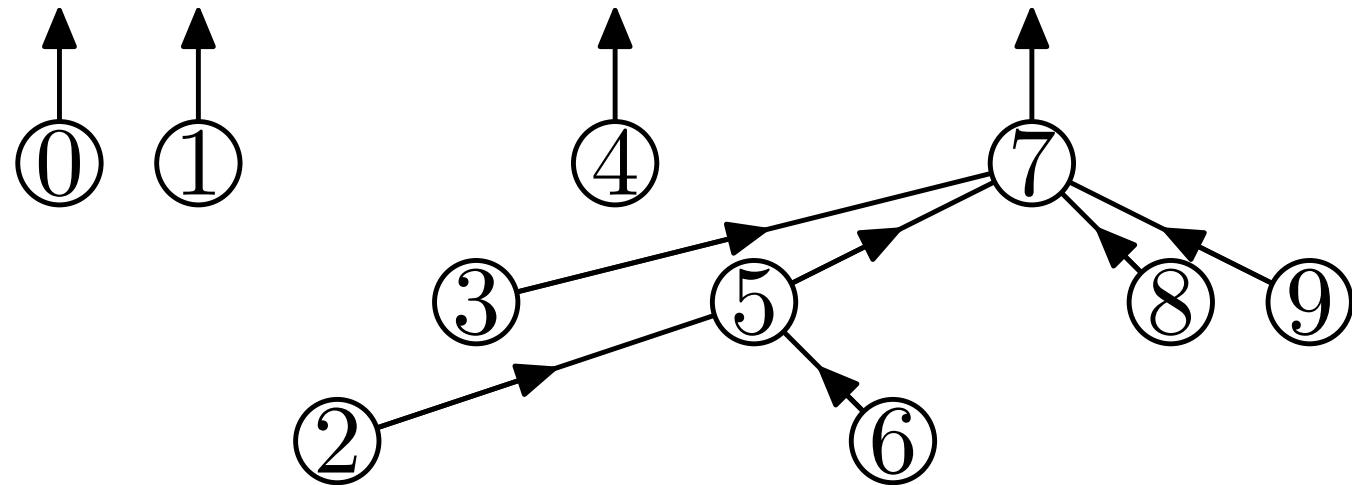
`find(3)`



-1	-1	5	5	-1	7	5	-3	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

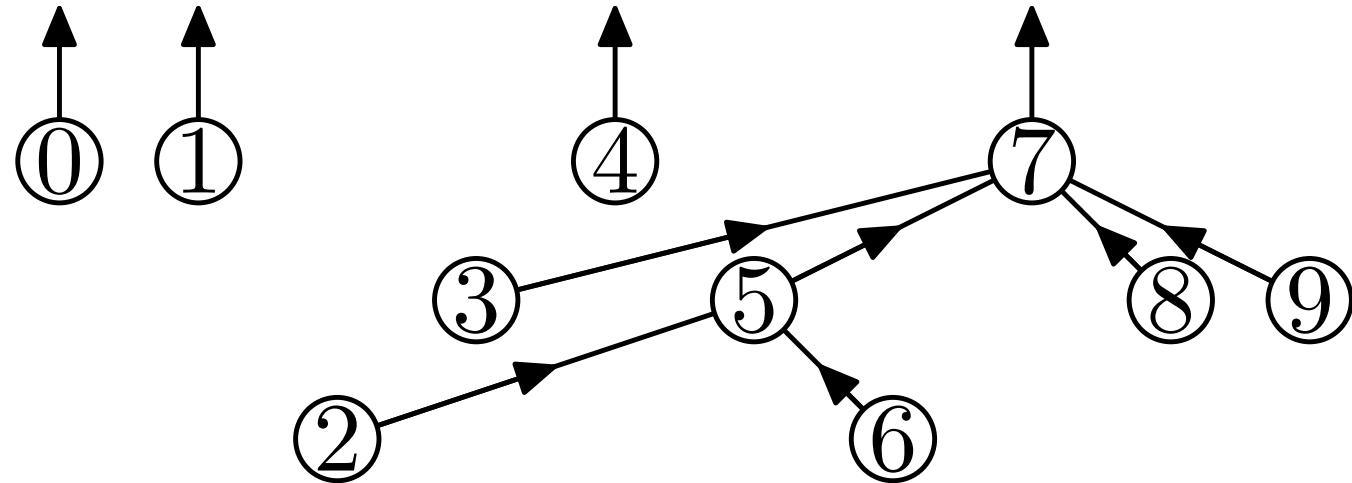
$\text{find}(3)=7$



-1	-1	5	7	-1	7	5	-3	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

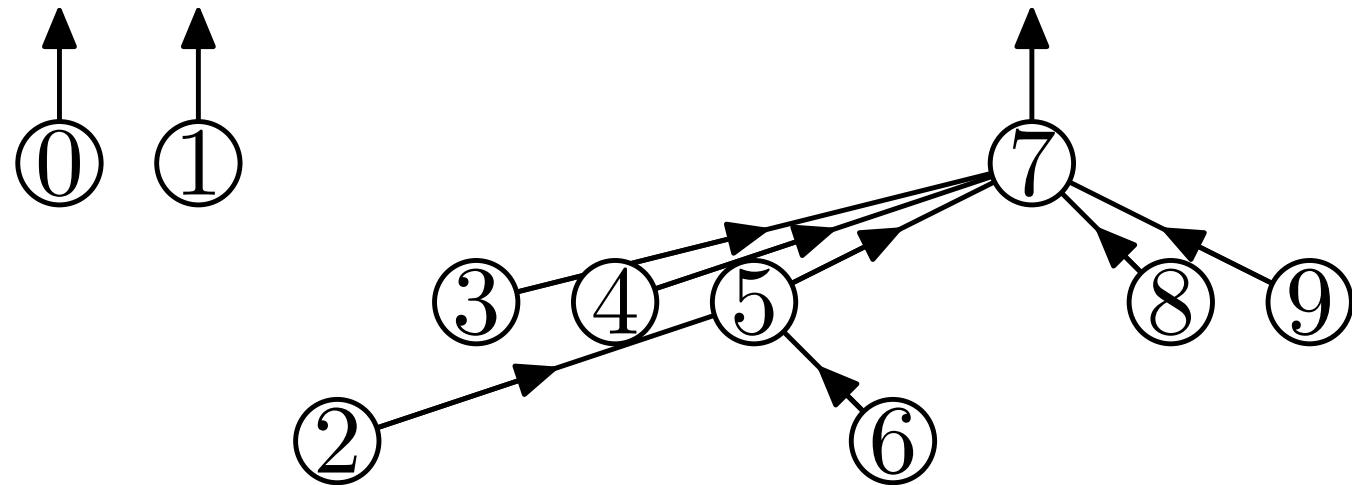
`union(find(3),find(4))`



-1	-1	5	7	-1	7	5	-3	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

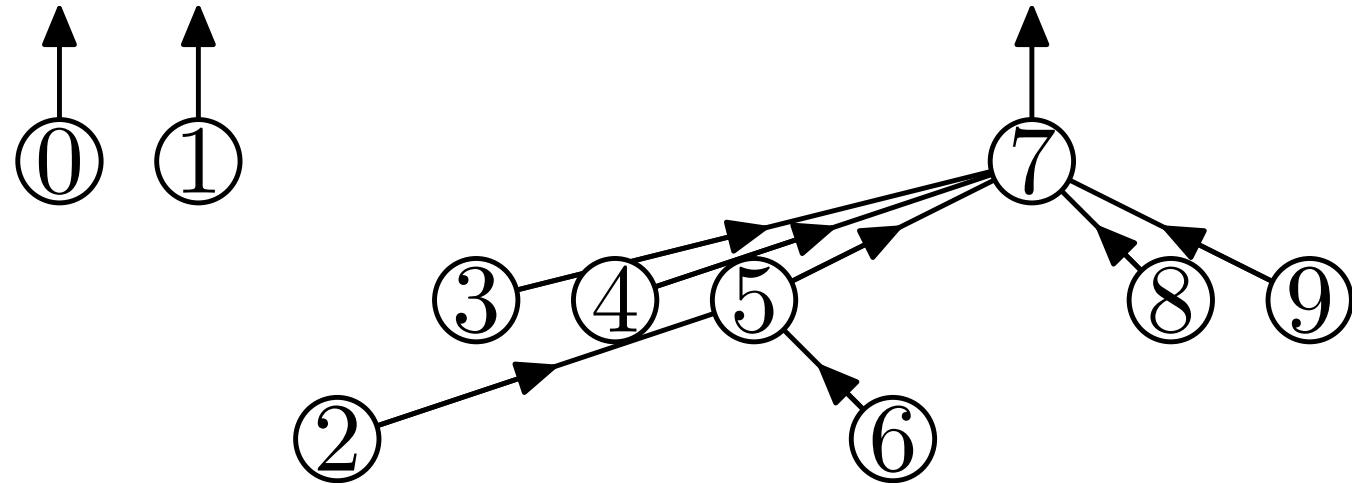
`union(find(3),find(4))`



-1	-1	5	7	7	7	5	-3	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

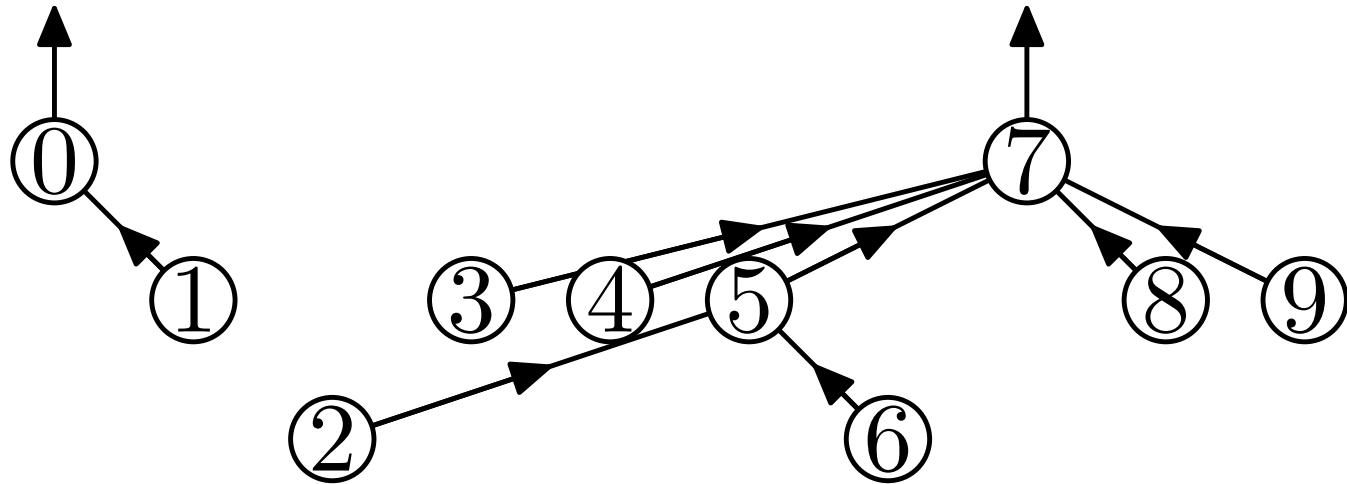
`union(find(0),find(1))`



-1	-1	5	7	7	7	5	-3	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

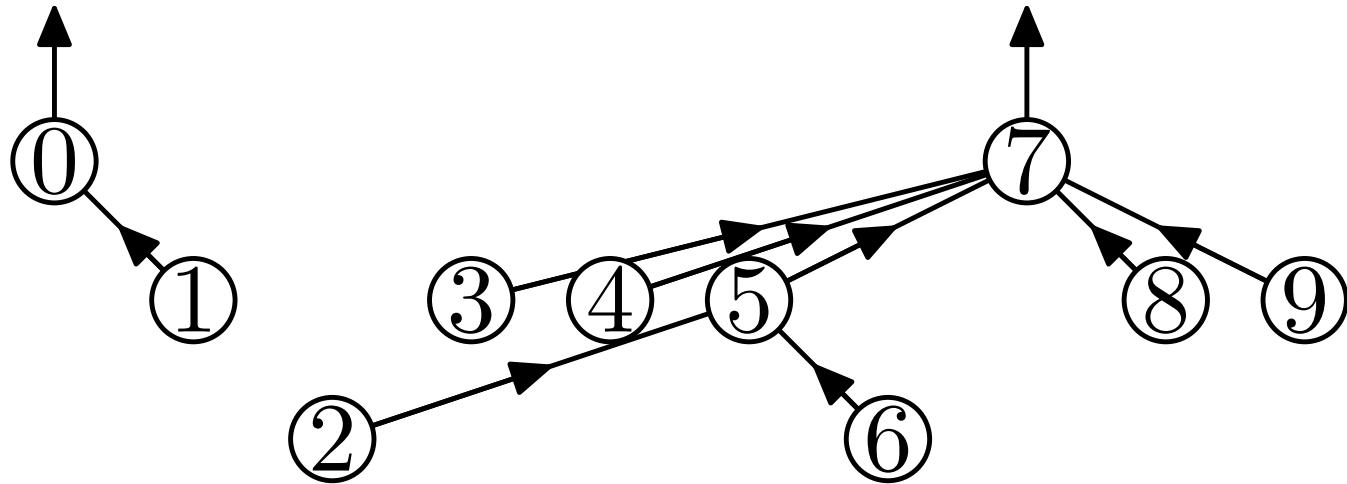
`union(find(0),find(1))`



-2	0	5	7	7	7	5	-3	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

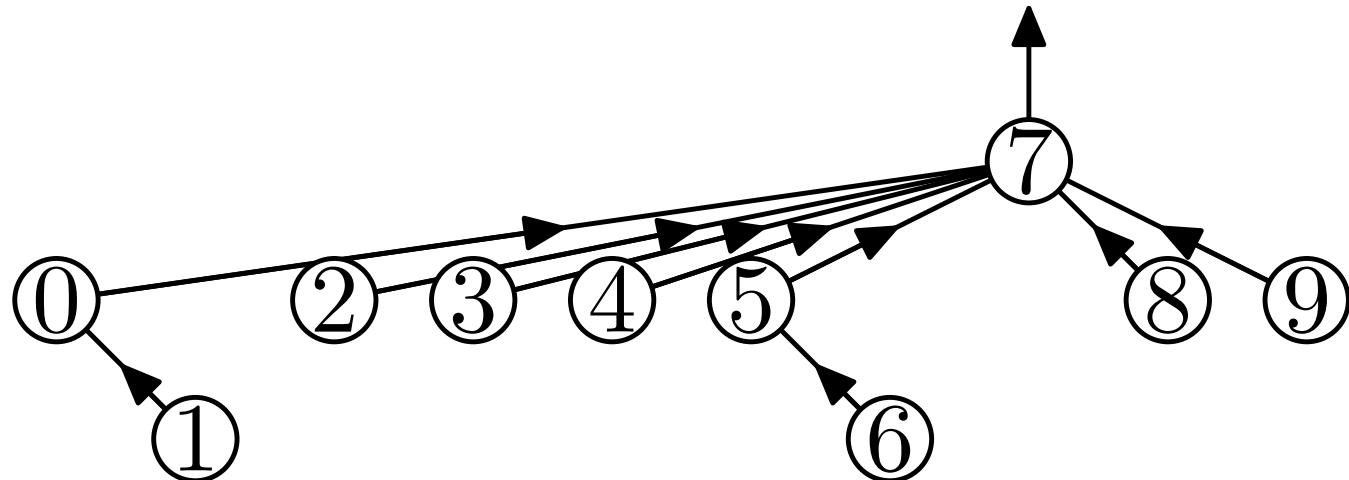
`union(find(1),find(2))`



-2	0	5	7	7	7	5	-3	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

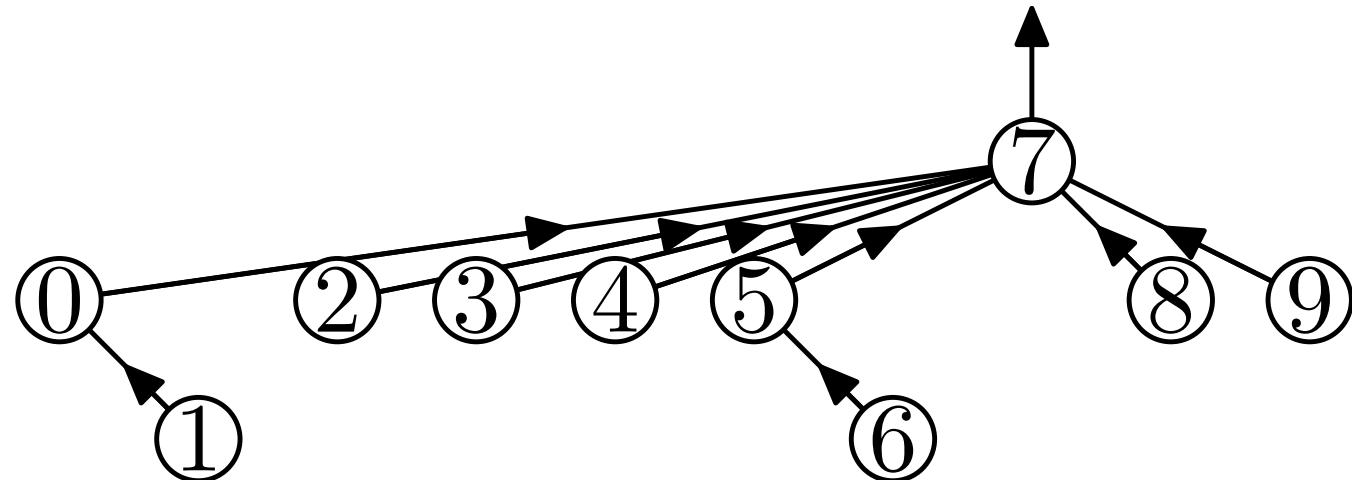
`union(find(1),find(2))`



7	0	7	7	7	7	5	-3	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

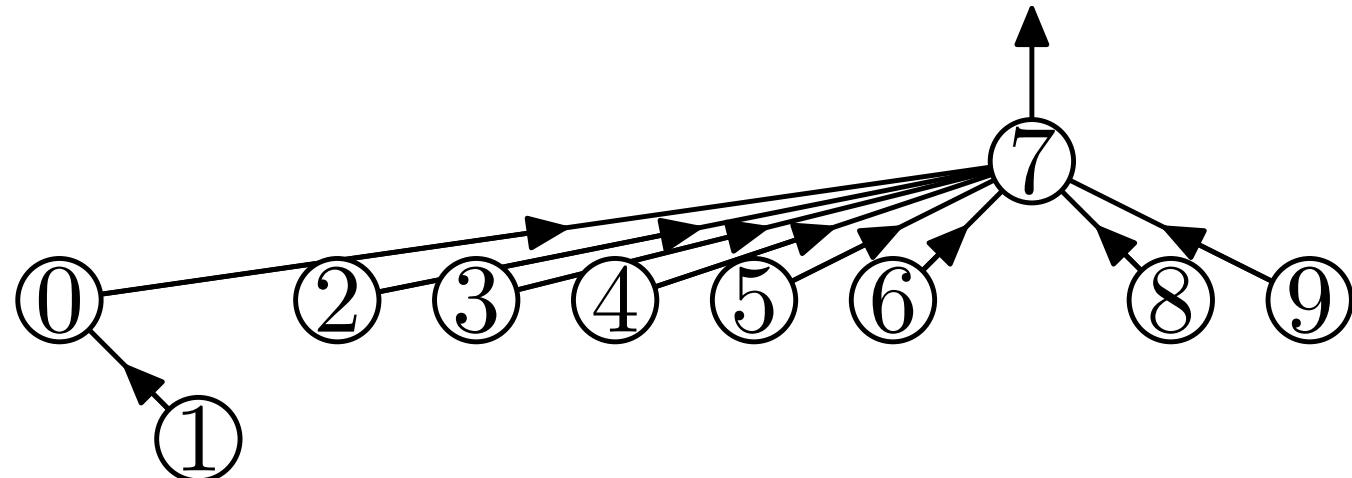
find(6)



7	0	7	7	7	7	5	-3	7	7
0	1	2	3	4	5	6	7	8	9

Putting it Together

$\text{find}(6)=7$

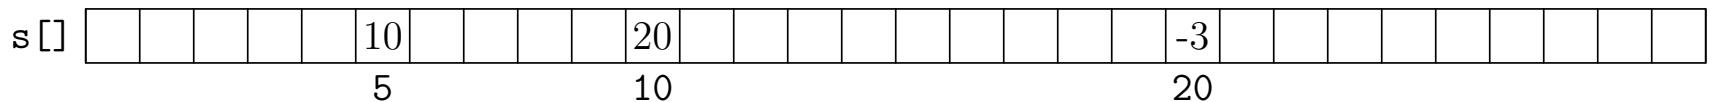


7	0	7	7	7	7	7	-3	7	7
0	1	2	3	4	5	6	7	8	9

Smart Union

```
DisjointSets::DisjointSets(int numElements)
{
    s = new int[numElements];
    for(int i=0; i<numElements; i++)
        s[i] = -1;                                // roots are negative number
}

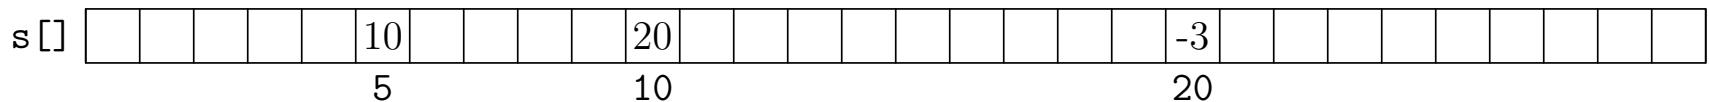
void DisjointSets::union_(int root1, int root2)
{
    if (s[root2]<s[root1]) {                  // root2 is deeper
        s[root1] = root2;                      // make root2 the root
    } else {
        if (s[root1]==s[root2])
            s[root1]--;                         // update height if same
        s[root2] = root1;                      // make root1 new root
    }
}
```



Smart Union

```
DisjointSets::DisjointSets(int numElements)
{
    s = new int[numElements];
    for(int i=0; i<numElements; i++)
        s[i] = -1;                                // roots are negative number
}

void DisjointSets::union_(int root1, int root2)
{
    if (s[root2]<s[root1]) {                  // root2 is deeper
        s[root1] = root2;                      // make root2 the root
    } else {
        if (s[root1]==s[root2])
            s[root1]--;                         // update height if same
        s[root2] = root1;                      // make root1 new root
    }
}
```



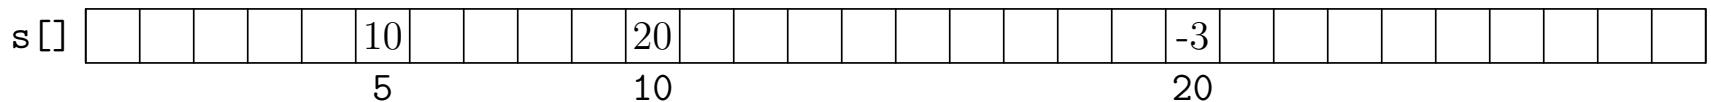
Smart Union

```

DisjointSets::DisjointSets(int numElements)
{
    s = new int[numElements];
    for(int i=0; i<numElements; i++)
        s[i] = -1;                                // roots are negative number
}

void DisjointSets::union_(int root1, int root2)
{
    if (s[root2]<s[root1]) {                  // root2 is deeper
        s[root1] = root2;                      // make root2 the root
    } else {
        if (s[root1]==s[root2])
            s[root1]--;                         // update height if same
        s[root2] = root1;                      // make root1 new root
    }
}

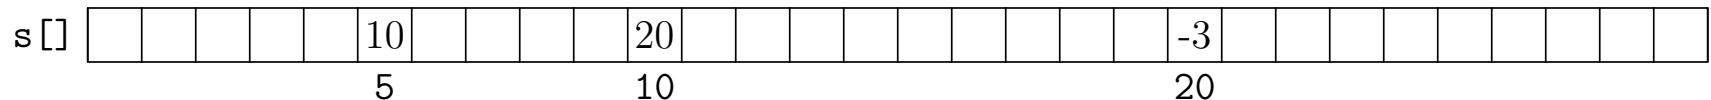
```



Path Compression

- To speed up `find` we relabel all nodes we visit during `find` by the root label

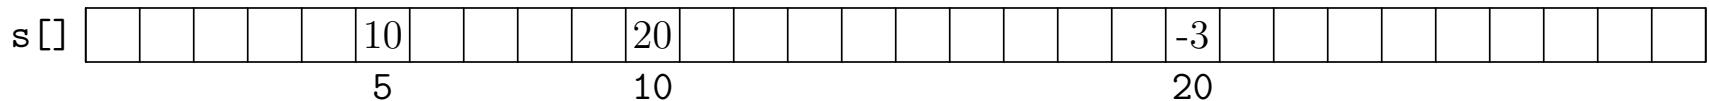
```
int DisjointSets::find(int index)
{
    if (s[index]<0)
        return index;
    else
        return s[index] = find(s[index]);
}
```



Path Compression

- To speed up `find` we relabel all nodes we visit during `find` by the root label

```
int DisjointSets::find(int index)
{
    if (s[index]<0)
        return index;
    else
        return s[index] = find(s[index]);
}
```



Mazes

- Union-Find is a data structure which can occur in very different applications
- One application is building a maze
- Start from a complete lattice
- Remove a randomly chosen edge if it connects two unconnected regions
- Stop when the start and end cell are connected
- Or better after all cells are connected

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24
25	26	27	28	29
30	31	32	33	34
35	36	37	38	39
40	41	42	43	44
45	46	47	48	49

Mazes

- Union-Find is a data structure which can occur in very different applications
- One application is building a maze
- Start from a complete lattice
- Remove a randomly chosen edge if it connects two unconnected regions
- Stop when the start and end cell are connected
- Or better after all cells are connected

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24
25	26	27	28	29
30	31	32	33	34
35	36	37	38	39
40	41	42	43	44
45	46	47	48	49

Mazes

- Union-Find is a data structure which can occur in very different applications
- One application is building a maze
- **Start from a complete lattice**
- Remove a randomly chosen edge if it connects two unconnected regions
- Stop when the start and end cell are connected
- Or better after all cells are connected

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24
25	26	27	28	29
30	31	32	33	34
35	36	37	38	39
40	41	42	43	44
45	46	47	48	49

Mazes

- Union-Find is a data structure which can occur in very different applications
- One application is building a maze
- Start from a complete lattice
- Remove a randomly chosen edge if it connects two unconnected regions
- Stop when the start and end cell are connected
- Or better after all cells are connected

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24
25	26	27	28	29
30	31	32	33	34
35	36	37	38	39
40	41	42	43	44
45	46	47	48	49

Mazes

- Union-Find is a data structure which can occur in very different applications
- One application is building a maze
- Start from a complete lattice
- Remove a randomly chosen edge if it connects two unconnected regions
- Stop when the start and end cell are connected
- Or better after all cells are connected

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24
25	26	27	28	29
30	31	32	33	34
35	36	37	38	39
40	41	42	43	44
45	46	47	48	49

Mazes

- Union-Find is a data structure which can occur in very different applications
- One application is building a maze
- Start from a complete lattice
- Remove a randomly chosen edge if it connects two unconnected regions
- Stop when the start and end cell are connected
- Or better after all cells are connected

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24
25	26	27	28	29
30	31	32	33	34
35	36	37	38	39
40	41	42	43	44
45	46	47	48	49

Time Complexity of Union-Find

- If we perform M finds and N unions then the time complexity is $O(M \log_2^*(N))$
- Where $\log_2^*(N)$ is the number of times you need to apply the logarithm function before you get a number less than 1
- In practice $\log_2^*(N) \leq 5$ for all conceivable N

Time Complexity of Union-Find

- If we perform M finds and N unions then the time complexity is $O(M \log_2^*(N))$
- Where $\log_2^*(N)$ is the number of times you need to apply the logarithm function before you get a number less than 1
- In practice $\log_2^*(N) \leq 5$ for all conceivable N

Time Complexity of Union-Find

- If we perform M finds and N unions then the time complexity is $O(M \log_2^*(N))$
- Where $\log_2^*(N)$ is the number of times you need to apply the logarithm function before you get a number less than 1
- In practice $\log_2^*(N) \leq 5$ for all conceivable N

Time Complexity of Union-Find

- If we perform M finds and N unions then the time complexity is $O(M \log_2^*(N))$
- Where $\log_2^*(N)$ is the number of times you need to apply the logarithm function before you get a number less than 1
- In practice $\log_2^*(N) \leq 5$ for all conceivable N

$$\log_2(10^{80}) = 265.75$$

Time Complexity of Union-Find

- If we perform M finds and N unions then the time complexity is $O(M \log_2^*(N))$
- Where $\log_2^*(N)$ is the number of times you need to apply the logarithm function before you get a number less than 1
- In practice $\log_2^*(N) \leq 5$ for all conceivable N

$$\log_2(\log_2(10^{80})) = 8.0539$$

Time Complexity of Union-Find

- If we perform M finds and N unions then the time complexity is $O(M \log_2^*(N))$
- Where $\log_2^*(N)$ is the number of times you need to apply the logarithm function before you get a number less than 1
- In practice $\log_2^*(N) \leq 5$ for all conceivable N

$$\log_2(\log_2(\log_2(10^{80}))) = 3.0097$$

Time Complexity of Union-Find

- If we perform M finds and N unions then the time complexity is $O(M \log^*(N))$
- Where $\log^*(N)$ is the number of times you need to apply the logarithm function before you get a number less than 1
- In practice $\log^*(N) \leq 5$ for all conceivable N

$$\log_2(\log_2(\log_2(\log_2(10^{80})))) = 1.5896$$

Time Complexity of Union-Find

- If we perform M finds and N unions then the time complexity is $O(M \log_2^*(N))$
- Where $\log_2^*(N)$ is the number of times you need to apply the logarithm function before you get a number less than 1
- In practice $\log_2^*(N) \leq 5$ for all conceivable N

$$\log_2(\log_2(\log_2(\log_2(\log_2(10^{80})))))) = 0.66868$$

Time Complexity of Union-Find

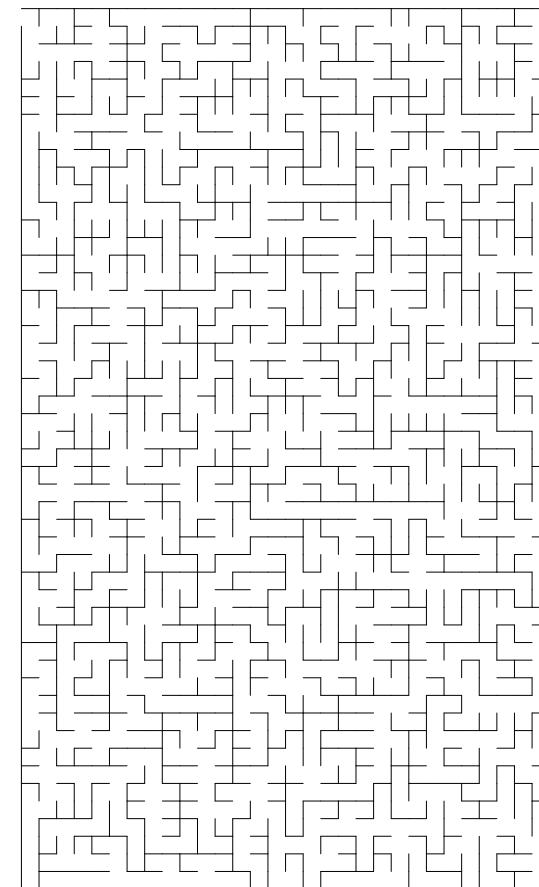
- If we perform M finds and N unions then the time complexity is $O(M \log^*(N))$
- Where $\log^*(N)$ is the number of times you need to apply the logarithm function before you get a number less than 1
- In practice $\log^*(N) \leq 5$ for all conceivable N

$$\log_2(\log_2(\log_2(\log_2(\log_2(10^{80})))))) = 0.66868$$

- The proof of this time complexity is rather involved

Outline

1. Equivalent Classes
2. Disjoint Sets
3. **Fast Sets**



Comparison of Sets

- Binary Search Trees: $O(\log_2(n))$, general purpose
- Hash tables: $O(1)$, but need to compute hash, slow iterator when sparse, general purpose
- B-trees: $O((k - 1) \log_k(n))$ very complicated, used for large amounts of data
- Tries: $O(\text{len}(\text{seq})) = O(\log_k(n))$ for large k expensive in memory, complicated to code efficiently

Comparison of Sets

- Binary Search Trees: $O(\log_2(n))$, general purpose
- Hash tables: $O(1)$, but need to compute hash, slow iterator when sparse, general purpose
- B-trees: $O((k - 1) \log_k(n))$ very complicated, used for large amounts of data
- Tries: $O(\text{len}(\text{seq})) = O(\log_k(n))$ for large k expensive in memory, complicated to code efficiently

Comparison of Sets

- Binary Search Trees: $O(\log_2(n))$, general purpose
- Hash tables: $O(1)$, but need to compute hash, slow iterator when sparse, general purpose
- B-trees: $O((k - 1) \log_k(n))$ very complicated, used for large amounts of data
- Tries: $O(\text{len}(\text{seq})) = O(\log_k(n))$ for large k expensive in memory, complicated to code efficiently

Comparison of Sets

- Binary Search Trees: $O(\log_2(n))$, general purpose
- Hash tables: $O(1)$, but need to compute hash, slow iterator when sparse, general purpose
- B-trees: $O((k - 1) \log_k(n))$ very complicated, used for large amounts of data
- Tries: $O(\text{len}(\text{seq})) = O(\log_k(n))$ for large k expensive in memory, complicated to code efficiently

What Set to Use?

- A PhD student and I were working on writing a fast solver for a combinatorial optimisation problem
- We had to choose one variable to change out of a small number of possible variables
- Each time we changed a variable then we had to update the list of possible variables (remove some variables add others)
- We wanted a data structure which had quick add and remove and where we could choose a variable at random

What Set to Use?

- A PhD student and I were working on writing a fast solver for a combinatorial optimisation problem
- We had to choose one variable to change out of a small number of possible variables
- Each time we changed a variable then we had to update the list of possible variables (remove some variables add others)
- We wanted a data structure which had quick add and remove and where we could choose a variable at random

What Set to Use?

- A PhD student and I were working on writing a fast solver for a combinatorial optimisation problem
- We had to choose one variable to change out of a small number of possible variables
- Each time we changed a variable then we had to update the list of possible variables (remove some variables add others)
- We wanted a data structure which had quick add and remove and where we could choose a variable at random

What Set to Use?

- A PhD student and I were working on writing a fast solver for a combinatorial optimisation problem
- We had to choose one variable to change out of a small number of possible variables
- Each time we changed a variable then we had to update the list of possible variables (remove some variables add others)
- We wanted a data structure which had quick add and remove and where we could choose a variable at random

What Set to Use?

- A PhD student and I were working on writing a fast solver for a combinatorial optimisation problem
- We had to choose one variable to change out of a small number of possible variables
- Each time we changed a variable then we had to update the list of possible variables (remove some variables add others)
- We wanted a data structure which had quick add and remove and where we could choose a variable at random—**what should we use?**

Bounded Set

- One special feature is that we knew we only wanted the set to contain integers between 0 and n (where n might be 100 000)
- This allowed us to use an array to represent whether an integer belongs to that set
- But how do we find a random element of the set quickly?

Bounded Set

- One special feature is that we knew we only wanted the set to contain integers between 0 and n (where n might be 100 000)
- This allowed us to use an array to represent whether an integer belongs to that set
- But how do we find a random element of the set quickly?

Bounded Set

- One special feature is that we knew we only wanted the set to contain integers between 0 and n (where n might be 100 000)
- This allowed us to use an array to represent whether an integer belongs to that set
- But how do we find a random element of the set quickly?

Bounded Set

- One special feature is that we knew we only wanted the set to contain integers between 0 and n (where n might be 100 000)
- This allowed us to use an array to represent whether an integer belongs to that set
- But how do we find a random element of the set quickly?
- Use another array of course!

FastSet

FastSet

0	1	2	3	4	5	6	7	8	9
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

FastSet

`add(4)`

0	1	2	3	4	5	6	7	8	9
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

FastSet

true

0	1	2	3	4	5	6	7	8	9
-1	-1	-1	-1	0	-1	-1	-1	-1	-1
4									

FastSet

`add(9)`

0	1	2	3	4	5	6	7	8	9
-1	-1	-1	-1	0	-1	-1	-1	-1	-1
4									

FastSet

true

0	1	2	3	4	5	6	7	8	9
-1	-1	-1	-1	0	-1	-1	-1	-1	1
4	9								

FastSet

`add(7)`

0	1	2	3	4	5	6	7	8	9
-1	-1	-1	-1	0	-1	-1	-1	-1	1
4	9								

FastSet

true

0	1	2	3	4	5	6	7	8	9
-1	-1	-1	-1	0	-1	-1	2	-1	1
4	9	7							

FastSet

`add(4)`

0	1	2	3	4	5	6	7	8	9
-1	-1	-1	-1	0	-1	-1	2	-1	1
4	9	7							

FastSet

false

0	1	2	3	4	5	6	7	8	9
-1	-1	-1	-1	0	-1	-1	2	-1	1
4	9	7							

FastSet

`add(1)`

0	1	2	3	4	5	6	7	8	9
-1	-1	-1	-1	0	-1	-1	2	-1	1
4	9	7							

FastSet

true

0	1	2	3	4	5	6	7	8	9
-1	3	-1	-1	0	-1	-1	2	-1	1
4	9	7	1						

FastSet

contains(9)

0	1	2	3	4	5	6	7	8	9
-1	3	-1	-1	0	-1	-1	2	-1	1
4	9	7	1						

FastSet

true

0	1	2	3	4	5	6	7	8	9
-1	3	-1	-1	0	-1	-1	2	-1	1
4	9	7	1						

FastSet

contains(5)

0	1	2	3	4	5	6	7	8	9
-1	3	-1	-1	0	-1	-1	2	-1	1
4	9	7	1						

FastSet

false

0	1	2	3	4	5	6	7	8	9
-1	3	-1	-1	0	-1	-1	2	-1	1
4	9	7	1						

FastSet

remove(9)

0	1	2	3	4	5	6	7	8	9
-1	3	-1	-1	0	-1	-1	2	-1	1
4	9	7	1						

FastSet

true

0	1	2	3	4	5	6	7	8	9
-1	1	-1	-1	0	-1	-1	2	-1	-1
4	1	7							

Implementation

```
class FastSet {  
    private:  
        int* indexArray;  
        int* memberArray;  
        int noMembers;  
  
    public:  
        FastSet (int n) {  
            indexArray = new int[n];  
            memberArray = new int[n];  
            for(int i=0; i<n; i++) {  
                indexArray [i] = -1;  
            }  
            noMembers = 0;  
        }  
  
        ~FastSet () {  
            delete[] indexArray;  
            delete[] memberArray;  
        }  
  
        int size () { return noMembers; }  
}
```

Implementation

```
class FastSet {  
    private:  
        int* indexArray;  
        int* memberArray;  
        int noMembers;  
  
    public:  
        FastSet(int n) {  
            indexArray = new int[n];  
            memberArray = new int[n];  
            for(int i=0; i<n; i++) {  
                indexArray[i] = -1;  
            }  
            noMembers = 0;  
        }  
  
        ~FastSet() {  
            delete[] indexArray;  
            delete[] memberArray;  
        }  
  
        int size() { return noMembers; }  
}
```

Implementation

```
class FastSet {  
    private:  
        int* indexArray;  
        int* memberArray;  
        int noMembers;  
  
    public:  
        FastSet (int n) {  
            indexArray = new int[n];  
            memberArray = new int[n];  
            for(int i=0; i<n; i++) {  
                indexArray [i] = -1;  
            }  
            noMembers = 0;  
        }  
  
        ~FastSet () {  
            delete[] indexArray;  
            delete[] memberArray;  
        }  
  
        int size() { return noMembers; }  
}
```

Add and Remove

```
bool add(int i) {
    if (indexArray[i]>-1)
        return false;
    memberArray [noMembers] = i;
    indexArray [i] = noMembers;
    ++noMembers;
    return true;
}

bool remove(int i) {
    if (indexArray[i]==-1)
        return false;
    --noMembers;
    memberArray [indexArray[i]] = memberArray [noMembers];
    indexArray [memberArray [noMembers]] = indexArray [i];
    indexArray [i] = -1;
    return true;
}
```

Add and Remove

```
bool add(int i) {
    if (indexArray[i]>-1)
        return false;
    memberArray [noMembers] = i;
    indexArray[i] = noMembers;
    ++noMembers;
    return true;
}

bool remove(int i) {
    if (indexArray[i]==-1)
        return false;
    --noMembers;
    memberArray [indexArray[i]] = memberArray [noMembers];
    indexArray [memberArray [noMembers]] = indexArray[i];
    indexArray[i] = -1;
    return true;
}
```

Collection Methods

```
void clear() {  
    for(int i=0; i<noMembers; i++) {  
        indexArray[memberArray[i]] = -1;  
    }  
    noMembers = 0;  
}  
  
bool isEmpty() {  
    return noMembers==0;  
}  
  
int* begin() {return &memberArray[0];}  
int* end() {return &memberArray[noMembers];}  
}
```

Collection Methods

```
void clear() {  
    for(int i=0; i<noMembers; i++) {  
        indexArray[memberArray[i]] = -1;  
    }  
    noMembers = 0;  
}  
  
bool isEmpty() {  
    return noMembers==0;  
}  
  
int* begin() {return &memberArray[0];}  
int* end() {return &memberArray[noMembers];}  
}
```

Collection Methods

```
void clear() {  
    for(int i=0; i<noMembers; i++) {  
        indexArray[memberArray[i]] = -1;  
    }  
    noMembers = 0;  
}  
  
bool isEmpty() {  
    return noMembers==0;  
}  
  
int* begin() {return &memberArray[0];}  
int* end() {return &memberArray[noMembers];}  
}
```

And Random?

- We can add additional methods taking advantage of the classes strength

private:

```
random_device rd; // Seed for the random number engine  
mt19937 gen(rd()); // Mersenne Twister RNG
```

public:

```
int getRandomElement () {  
    return memberArray[uniform_int_distribution<int>(0, noMembers)];  
}
```

- Need to use FastSet signature to use this

```
FastSet fastSet(n);  
:  
int r = fastSet.getRandomElement();
```

And Random?

- We can add additional methods taking advantage of the classes strength

private:

```
random_device rd; // Seed for the random number engine  
mt19937 gen(rd()); // Mersenne Twister RNG
```

public:

```
int getRandomElement () {  
    return memberArray[uniform_int_distribution<int>(0, noMembers)];  
}
```

- Need to use FastSet signature to use this

```
FastSet fastSet(n);  
:  
int r = fastSet.getRandomElement();
```

Speed Up

- We compared our algorithm to a very highly regarded “state-of-the-art” algorithm
- For large problems we were over 10 times faster because of this data structure
- The competitor algorithm used a complex tree structure instead of the simple array
- Why?

Speed Up

- We compared our algorithm to a very highly regarded “state-of-the-art” algorithm
- For large problems we were over 10 times faster because of this data structure
- The competitor algorithm used a complex tree structure instead of the simple array
- Why?

Speed Up

- We compared our algorithm to a very highly regarded “state-of-the-art” algorithm
- For large problems we were over 10 times faster because of this data structure
- The competitor algorithm used a complex tree structure instead of the simple array
- Why?

Speed Up

- We compared our algorithm to a very highly regarded “state-of-the-art” algorithm
- For large problems we were over 10 times faster because of this data structure
- The competitor algorithm used a complex tree structure instead of the simple array
- Why?

Speed Up

- We compared our algorithm to a very highly regarded “state-of-the-art” algorithm
- For large problems we were over 10 times faster because of this data structure
- The competitor algorithm used a complex tree structure instead of the simple array
- Why? **The array solution isn't in the books**

Lessons

- If you have a bounded set then using an array is usually going to be very fast $O(1)$ (or $O(\log^*(n))$)
- These data structures are not general purpose for solving every day problems (c.f. `vector<T>`, `set<T>` and `map<T>`)
- They are “back pocket” data structures that solve problems that come up often enough that they are worth knowing about
- Sometimes good algorithms are not documented, but it doesn’t mean they don’t exist

Lessons

- If you have a bounded set then using an array is usually going to be very fast $O(1)$ (or $O(\log^*(n))$)
- These data structures are not general purpose for solving everyday problems (c.f. `vector<T>`, `set<T>` and `map<T>`)
- They are “back pocket” data structures that solve problems that come up often enough that they are worth knowing about
- Sometimes good algorithms are not documented, but it doesn’t mean they don’t exist

Lessons

- If you have a bounded set then using an array is usually going to be very fast $O(1)$ (or $O(\log^*(n))$)
- These data structures are not general purpose for solving every day problems (c.f. `vector<T>`, `set<T>` and `map<T>`)
- They are “back pocket” data structures that solve problems that come up often enough that they are worth knowing about
- Sometimes good algorithms are not documented, but it doesn’t mean they don’t exist

Lessons

- If you have a bounded set then using an array is usually going to be very fast $O(1)$ (or $O(\log^*(n))$)
- These data structures are not general purpose for solving every day problems (c.f. `vector<T>`, `set<T>` and `map<T>`)
- They are “back pocket” data structures that solve problems that come up often enough that they are worth knowing about
- Sometimes good algorithms are not documented, but it doesn’t mean they don’t exist