

Lesson 11: Use Heaps!



Heaps, Priority queues, Heap Sort, Other heaps


AICE1005

Algorithms and Analysis

1

Heaps

- A (min-)heap is (from one perspective) a binary tree
- It is a binary tree satisfying two constraints
 - It is a **complete** tree
 - Each child has a value 'greater than or equal to' its parent



- complete** means that every level is fully occupied above the lowest level and the nodes on the lowest level are all to the left

AICE1005

Algorithms and Analysis

3

Priority Queues

- One of the prime uses of heaps is to implement a priority queue
- A Priority Queue is a queue with priorities
- That is, we assign a priority to each element we add
- The head of the queue is the element with highest priority (smallest number)
- Used, for example, in simulating real time events
- Used to implement "greedy algorithms"


AICE1005

Algorithms and Analysis

5

Array Implementation of Heaps

- Because the tree is complete we can implement the heap efficiently using an array



AICE1005

Algorithms and Analysis

7

- Heaps
- Priority Queues
 - Array Implementation
- Heap Sort
- Other Heaps



AICE1005

Algorithms and Analysis

2

Outline

- Heaps
- Priority Queues
 - Array Implementation
- Heap Sort
- Other Heaps



AICE1005

Algorithms and Analysis

4

Priority Queue

- A simple Priority Queue might include
 - unsigned size() returning the the number of elements
 - bool empty() returns true if empty
 - void push(T element, int priority) adds an element
 - T top() returns head of queue
 - void pop() dequeues head of queue

AICE1005

Algorithms and Analysis

6

Code for a Priority Queue

```
#include <vector>
using namespace std;

template <typename T, typename P>
class heapPQ {
private:
    vector<pair<T, P> > array;

public:
    heapPQ(unsigned capacity=11) {
        array.reserve(capacity);
    }

    unsigned size() {return array.size();}

    bool empty() {return array.empty();}


    const T& top() {return array[0].first;}
```

AICE1005

Algorithms and Analysis

8

Adding an Element



AIICE1005

Algorithms and Analysis

9

Adding an Element

```
void push(T value, P priority) {
    pair<T,P> tmp(value, priority);
    array.push_back(tmp);
    unsigned child = size() - 1;


    /* Percolate Up */
    while(child!=0) {
        unsigned parent = (child-1)>>1; // floor((child-1)/2)
        if (array[parent].second < array[child].second)
            return;
        array[child] = array[parent];
        array[parent] = tmp;
        child = parent;
    }
}
```

AIICE1005

Algorithms and Analysis

10

Popping the Top



AIICE1005

Algorithms and Analysis

11

Popping the Top

```
void pop() {
    unsigned parent = 0;
    pair<T, P> tmp = array.back();
    array[0] = tmp;
    array.pop_back();
    unsigned child = 1;


    /* Percolate down */
    while(child<size()) {
        if (child+1<size() && array[child+1].second < array[child].second)
            ++child;
        if (array[child].second > array[parent].second)
            return;
        array[parent] = array[child];
        array[child] = tmp;
        parent = child;
        child = 2*parent + 1;
    }
}
```

AIICE1005

Algorithms and Analysis

12

Heaps in Action



AIICE1005

Algorithms and Analysis

13

Time Complexity of Heaps

- The two important operation are add and removeMin
- These both either percolating an element up the tree or percolating an element down the tree
- The number of operations depends on the depth of the tree which is $\Theta(\log(n))$
- Thus add and removeMin are $O(\log(n))$
- Except add could also require resizing the array, but the amortised cost of this is low


AIICE1005

Algorithms and Analysis

14

Real Time Simulation

- A nice application of priority queues is to perform real time simulations
- I was once modelling a neural network where neuron fired an impulse which would then be received by other neurons



AIICE1005

Algorithms and Analysis

15

Synchronised Firing

- We wanted to show that if a group of neurons fired together they could make another group of neurons fire in synchrony despite the fact that it would take different times for the receiving neurons to feel the pulse (due to the different lengths of the axons)
- A famous Israeli group had "proved" this couldn't happen
- Using a priority queue we modelled the neurons
 - When a neuron fired the receiving neurons would be put on a priority queue according to when they received the pulse
 - If the receiving neurons received enough pulse in a short enough time they would then fire

AIICE1005

Algorithms and Analysis

16

- Using a priority queue meant we knew when the next event would happen
- We did not have to run a clock where most of the time nothing happened
- This allowed us to perform a very large simulation efficiently
- The simulation showed that the pulse of neurons synchronised despite the “proof” that this wouldn’t happen

Heap Sort

- A priority queue suggests a very simple way of performing sort
- We simply add elements to a heap and then take them off again


```
template <typename T>
void sort(vector<T> aList)
{
    HeapPQ<T> aHeap = new HeapPQ<T>(aList.size());
    for (T element: aList)
        aHeap.push(element, element);

    aList.clear();
    while(aHeap.size() > 0) {
        aList.push_back(aHeap.top());
        aHeap.pop();
    }
}
```

- Note that this is not an in-place sort algorithm (i.e. it uses lots of memory)

Complexity of Heap Sort

- As we have to add n elements and then remove n elements the time complexity is log-linear, i.e. $O(n \log(n))$



- This is actually a very efficient algorithm


Other Heaps

- Binary Heaps are so useful that other types of heaps have been developed
- The simplest enhancement is to combine a binary heap with a map which maintains a pointer to each element
- The map has to be updated every time elements are moved in the heap (fortunately only $O(\log(n))$ elements are move each time the heap is updated)
- The advantage of this heap is that the priorities of elements can be changed (involving percolating elements up or down the tree)

1. Heaps
2. Priority Queues
 - Array Implementation
3. Heap Sort
4. Other Heaps



Example of Heap Sort



Outline

1. Heaps
2. Priority Queues
 - Array Implementation
3. Heap Sort
4. Other Heaps



Merging Heaps

- One common demand on a heaps is to merge two heaps
- Unfortunately binary heaps are not efficiently merged
- There are a variety of different heaps (leftist heaps, skew heaps, binomial queues, . . .) designed to be merged
- All these heaps are real binary trees (i.e. represented by pointers rather than put in an array)
- They are slower than a binary heap because indexing is slightly slower as is creating node objects
- However, they allow merging

- All other operations are achieved by merging
 - ★ Adding an element is achieved by merging the current heap with a heap of one element
 - ★ Removing the minimum element is achieved by removing the root and merging the left and right tree
- For details see the course text (you won't be examined on the details of these heaps)

- Heaps are a powerful data structure—they are particularly useful for implementing priority queues
- Heaps are binary trees that can be implemented as arrays
- Priority queue have many (often surprising uses)
 - ★ They are used when you need a queue with priorities, e.g. in operating systems
 - ★ They can be used to perform pretty efficient sort
 - ★ They are often used for implementing greedy type algorithms
 - ★ One important application is in real time simulations
- There exists many extensions of heaps