

Algorithms and Analysis

Lesson 25: *Settle For Good Solutions*



neighbourhood search, heuristics, simulated annealing, GA

Outline

1. Heuristic Search

- Constructive algorithms
- Neighbourhood search

2. Simulated Annealing

3. Evolutionary Algorithms



Heuristic Algorithms


- Given that we know of no efficient algorithms for finding the optimal solution to NP-hard problems we must content ourselves with either
 - ★ Spending a very long time (e.g. using branch and bound)
 - ★ Accepting good solutions which aren't necessarily optimal
- Algorithms for finding good solutions are often called **approximation algorithms** or **heuristic algorithms**

Heuristics

- The idea behind heuristic algorithms is to use a rough guide or **heuristic** pointing you in reasonable direction■
- If this heuristic is good you should find good solutions much faster than exhaustive search■
- Two commonly used heuristics are
 - ★ A greedy heuristic (take the best move)■
 - ★ Believe that good solutions are 'close' to each other■

Constructive Algorithms

- Constructive algorithms build-up a solution
- They usually rely on a greedy heuristic
- They are very fast
- Once you have got a solution that's it
- They can give reasonable solutions quickly, but they are not usually very good



Neighbourhood search


- An alternative to constructive algorithms are search algorithms relying on good solutions being close to each other■
- In neighbourhood search we■
 1. Start from some solution■
 2. Examine the neighbouring solutions ■
 3. Move to a neighbour if it is better or, at least, not worse■
 4. Repeat 2 until some stopping criteria■
- If we are maximising this is often called a **hill-climber**■
- If we are minimising it is often called **descent**■

Iterative Improvement at its Best

- There are times when a neighbourhood algorithm will find the optimal solution■
- The classic example of this is in **linear programming** where the simplex method leads to the optimal solution■
- Other examples include
 - ★ Maximum Flow
 - ★ Maximum Matching in Bipartite Graphs■
- Unfortunately, this doesn't always work■

Local Optima

- Neighbourhood search is usually much slower than a constructive algorithm but tends to find better quality solutions■
- However, it will often get stuck



Simple Fixes

- One simple fix is to restart from many different starting positions■
- Or perturb the current solution and restart■
- These give improvements over doing nothing, but aren't necessarily great strategies■
- You can also increase the size of the neighbour to decrease the chance of getting stuck (e.g. in TSP swap more cities)■

Outline

1. Heuristic Search

- Constructive algorithms
- Neighbourhood search

2. Simulated Annealing

3. Evolutionary Algorithms




Simulated Annealing

- Simulated annealing is an example of a stochastic hill-climber■
- Sometimes you go in the wrong direction (down-hill)■
- Historically it is an idea from physics■—where you tend to minimise energy■
- Idea is to obtain a (low energy) crystalline material you very slowly let the material cool from a liquid state (opposite of quenching)■

Stochastic Descent

- It is easier to fall down hill than to go back up



Simulated Annealing

- Algorithm to minimise energy $E(\mathbf{X})$ where $\mathbf{X} = (X_1, X_2, \dots, X_N)$ ■
- Starting from a random configuration \mathbf{X} ■
- Choose a neighbour \mathbf{X}' ■
- If the neighbour is better (lower energy) move to it■
- Otherwise move to the neighbour with some probability■
- The parameter β controls the probability of moving to a neighbour■
- We increase β to reduce the probability of going uphill over time■

Cooling Schedule

- The parameter β is known as the inverse temperature because of an analogy with physics■
- Over time we have to increase β (decrease the temperature) so that the system will remain in the low energy state■
- The way you reduce the temperature (increase β) is known as the cooling schedule■
- Choosing a good cooling schedule can be critical■
- Choosing a good cooling schedule is something of a black art■

Convergence Theorem

- There is a theorem that says if you choose a slow enough cooling schedule you will end up in the global minimum eventually■
- Unfortunately eventually is a very long time■
- It is quicker to search all possible states■
- Still people get very excited about convergence proofs■

Outline

1. Heuristic Search
 - Constructive algorithms
 - Neighbourhood search
2. Simulated Annealing
3. **Evolutionary Algorithms**




Genetic Algorithms

- Genetic algorithms are methods to evolve a population of potential candidates to find a good solution to an optimisation problem■
- There are a whole set of related methods that go by the name of evolutionary algorithms, GAs are a subspecies of EAs■
- They can be viewed as an engineering approach to solving hard problems■
- I'm going to present my, highly prejudiced view of what's important in making a GA work■

A Canonical GA

1. *Initialise population*■
2. **for** $t = 1$ **to** T ■
 - (a) *Evaluate fitness*■
 - (b) *Select* a new population based on fitness■
 - (c) *Mutate* members of the population■
 - (d) *Crossover* members of the population■
3. Return best member of the population■

E.g. Graph Colouring




- Given a graph $G = (\mathcal{V}, \mathcal{E})$ ■
- Assign colours, $c(v)$, to the vertices of the graph $v \in \mathcal{V}$ ■
- Minimise the number of edges $e = (v, v') \in \mathcal{E}$ with the same coloured vertices $c(v) = c(v')$ (colour conflicts) ■


Initialise Population

Generate random colourings E.g.


Generation 0: evaluate fitness




Cost = 12




Cost = 16




Cost = 7



Cost = 11



Cost = 13



Cost = 16

Selection


- Select a new population of P members preferentially choosing the fitter members■
- Let w_α be a measure of the fitness of member α ■
- E.g. choose members α with a probability

$$p_\alpha = \frac{w_\alpha}{\sum_{\alpha'=1}^P w_{\alpha'}} \quad \blacksquare$$

- Many different ways of doing this■(some better than others)■


Mutation

Change the colour of one or more of the vertices. E.g.



Crossover

Take two solutions and combine them to form a new solution. E.g.



Crossover Operators

- Single-point crossover

- ★ Take two strings and cut them at some random site

$$\left(\begin{array}{c|c} GRBGBR & BGGBGBG \end{array} \right) \rightarrow \left\{ \begin{array}{c|c} GRBGBR & RGRBBGB \\ RRBRGB & BGGBGBG \end{array} \right\}$$

- Multi-point crossover

- ★ Take two strings and cut them at several sites

$$\left(\begin{array}{c|c|c} GRBGBR & BGGB & GBG \\ RRBRGB & RGRB & BGB \end{array} \right) \rightarrow \left\{ \begin{array}{c|c|c} GRBGBR & RGRB & GBG \\ RRBRGB & BGGB & BGB \end{array} \right\}$$





- Uniform Crossover

- ★ Take two strings and create children by a random shuffle

$$\left(\begin{array}{c} GRBGBR RBGGBGBG \\ RRBRGB RGRBBGB \end{array} \right) \rightarrow \left\{ \begin{array}{c} GRBRGB RGRGGGB \\ RRBGBR RBGGBBBG \end{array} \right\}$$


- Any of these crossover can be biased towards one parent

Cost of Crossover




GA


Final Population




Cost = 8




Cost = 10




Cost = 7



Cost = 8



Cost = 8



Cost = 12

Bit Simulated Crossover

- Choose each variable independently with the probability proportional to the frequency of the allele in the population■

(\dots, B, \dots) (\dots, R, \dots) (\dots, G, \dots)
 (\dots, R, \dots) (\dots, B, \dots) (\dots, G, \dots) ■
 (\dots, B, \dots) (\dots, G, \dots) (\dots, B, \dots)
 (\dots, B, \dots)

$$p_i(B) = 0.5, \quad p_i(G) = 0.3, \quad p_i(R) = 0.2$$

- New algorithms built on this idea, “Estimation of Distribution Algorithms” EDAs■




Galinier and Hao's Crossover Operator

- Choose two parents■
- Sort nodes into colour-groups■

	Parent 1	Parent 2
B	{1,3,4,8,... }	{3,5,7,10, ... }
G	{2,6,7,10,... }	{1,11,12,13, ... }
R	{5,9,11,12,... }	{2,3,6,8, ... }
⋮	⋮	⋮

- Choose largest colour-group in parent 1■
- Eliminate all nodes from that colour-group in parent 2■
- Choose largest colour-group in parent 2■
- etc.■

Cost of Crossover



Other Heuristics

- There are many extensions to neighbourhood search, simulated annealing and genetic algorithms■
- There are other techniques such as Tabu search
 - ★ Construct a list of place you cannot go to (usually the last few configurations)■
 - ★ Make the best move you are allowed to make■
 - ★ Rather a large number of *ad hoc* rules to make it work■
 - ★ Often very fast but runs out of steam■
- Many other EAs including particle swarm optimisations (PSO), ant colony optimisation (ACO), evolutionary strategies, . . . ■

Which Heuristic is Best?

- The best heuristic depends on the application■
- Descent is very fast, but only finds local optima—good starting place■
- Tabu search is often very fast, but sometimes fails to find really good solutions■
- Simulated annealing and Genetic Algorithms are slow, but often find good solutions■
- The best algorithms tend to be special purpose algorithms designed for the problem■

Lessons

- For many problems the best strategy is to find a good solution, but not the best■
- Iterative search usually give good quality solutions■
- There are many variants of heuristic search■
- Heuristic search algorithms aren't fast (don't use these techniques in an interactive program if you want to keep customers)■
- For large combinatorial optimisation problems this is often the only choice■